

Importing relevant libraries.

```
In [2]: import numpy as np
import pandas as pd
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        os.path.join(dirname, filename)
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
import cv2 as cv
import numpy as np
from matplotlib import pyplot as plt
import pandas as pd
import pydicom
from skimage.transform import resize
import matplotlib.patches as patches
from tqdm import tqdm
import math

import tensorflow as tf
from tensorflow.keras.utils import plot_model
```

Loading labels and checking dimensions.

```
In [4]: train_label = pd.read_csv(r'C:\Users\Sanan\Downloads\CompApp Project\rsna-pneumonia
train_label.shape
```

```
Out[4]: (30227, 6)
```

Creating function designed to resize an image to fit within a fixed input\_size while maintaining aspect ratio, while adjusting the bounding box accordingly.

```
In [5]: input_size = 244

def format_image(img, box):
    height, width = img.shape
    max_size = max(height, width)
    r = max_size / input_size
    new_width = int(width / r)
    new_height = int(height / r)
    new_size = (new_width, new_height)
    resized = cv.resize(img, new_size, interpolation=cv.INTER_LINEAR)
    new_image = np.zeros((input_size, input_size), dtype=np.uint8)
    new_image[0:new_height, 0:new_width] = resized

    x, y, w, h = (box[0], box[1], box[2], box[3]) if box[0] else (0.0,0.0,0.0,0.0)
    new_box = [int((x)/ r), int((y)/ r), int(w/ r), int(h/ r)] if box[0] else [0.0,

    return new_image, new_box
```

This function processes a DICOM medical image, resizes it, adjusts the bounding box, converts it to RGB, and visualizes it with the bounding box overlay.


```
In [6]: datapath = 'rsna-pneumonia-detection-challenge\stage_2_train_images\00436515-870c-4
temp_img = pydicom.dcmread(datapath).pixel_array
temp_box = [264.0, 152.0, 213.0, 379.0]

temp_img_formated, box = format_image(temp_img, temp_box)
print(box)
temp_color_img = cv.cvtColor(temp_img_formated, cv.COLOR_GRAY2RGB)

cv.rectangle(temp_color_img, box, (0, 255, 0), 1)

plt.imshow(temp_color_img)
# plt.axis("off")
plt.show()
```

[62, 36, 50, 90]



Load labels and display first 5 rows.

```
In [7]: train_labels = pd.read_csv('rsna-pneumonia-detection-challenge\stage_2_train_labels
train_labels.head()
```

Out[7]:

	patientId	x	y	width	height	Target
0	0004cfab-14fd-4e49-80ba-63a80b6bddd6	NaN	NaN	NaN	NaN	0
1	00313ee0-9eaa-42f4-b0ab-c148ed3241cd	NaN	NaN	NaN	NaN	0
2	00322d4d-1c29-4943-afc9-b6754be640eb	NaN	NaN	NaN	NaN	0
3	003d8fa0-6bf1-40ed-b54c-ac657f8495c5	NaN	NaN	NaN	NaN	0
4	00436515-870c-4b36-a041-de91049b9ab4	264.0	152.0	213.0	379.0	1

Creates and uses a function to load and process the DICOM images of Lung CTs and labels.

```
In [8]: def data_load(dataset, batch_size=3, full_data_path=r"rsna-pneumonia-detection-challenge",
X = []
Y = []

for index, row in tqdm(dataset.iterrows(), total=len(dataset), desc="Loading data"):
    filename = row['patientId']

    temp_img = pydicom.dcmread(os.path.join(full_data_path, filename + ".dcm"))
    temp_box = [row['x'], row['y'], row['width'], row['height']] if not math.isnan(row['x']) and not math.isnan(row['y']) and not math.isnan(row['width']) and not math.isnan(row['height']) else [None, None, None, None]

    img, box = format_image(temp_img, temp_box)

    img = img.astype(float) / 255.
    box = np.asarray(box, dtype=float) / input_size

    label = np.append(box, row['Target'])

    X.append(img)
    Y.append(label)

X = np.array(X)
data_X_len = len(X)
X = np.expand_dims(X, axis=3)
X = tf.convert_to_tensor(X, dtype=tf.float32)
Y = tf.convert_to_tensor(Y, dtype=tf.float32)

result = tf.data.Dataset.from_tensor_slices((X, Y))

return result, data_X_len

raw_train_ds, train_len = data_load(train_labels[:5200], ds_type="train")
print(train_len)
raw_valid_ds, valid_len = data_load(train_labels[5200:5900], ds_type="not train")
raw_test_ds, test_len = data_load(train_labels[5900:6501], ds_type="not train")
```

Loading data: 100%|██████████| 5200/5200 [02:41<00:00, 32.30it/s]  
5200

Loading data: 100%|██████████| 700/700 [00:23<00:00, 30.31it/s]  
Loading data: 100%|██████████| 601/601 [00:19<00:00, 30.72it/s]

Defines a function to ready the images, label pair to be used in TensorFlow.

```
In [9]: CLASSES = 2

def format_instance(image, label):
    return image, (tf.one_hot(int(label[4]), CLASSES), [label[0], label[1], label[2]
```

Defines a function to optimize the training dataset for the tensor flow model.

```
In [11]: BATCH_SIZE = 32

def tune_training_ds(dataset):
    dataset = dataset.map(format_instance, num_parallel_calls=tf.data.AUTOTUNE)
    dataset = dataset.shuffle(1024, reshuffle_each_iteration=True)
    dataset = dataset.repeat() # The dataset be repeated indefinitely.
    dataset = dataset.batch(BATCH_SIZE)
    dataset = dataset.prefetch(tf.data.AUTOTUNE)
    return dataset
```

Defines a function to optimize the validation dataset for the tensor flow model.

```
In [12]: def tune_validation_ds(dataset):
    dataset = dataset.map(format_instance, num_parallel_calls=tf.data.AUTOTUNE)
    dataset = dataset.batch(len(dataset) // 4)
    dataset = dataset.repeat()
    return dataset
```

Using the above functions, preparing the training and validation datasets.

```
In [13]: train_ds = tune_training_ds(raw_train_ds)
validation_ds = tune_validation_ds(raw_valid_ds)
```

Visualizing part of the training data set.

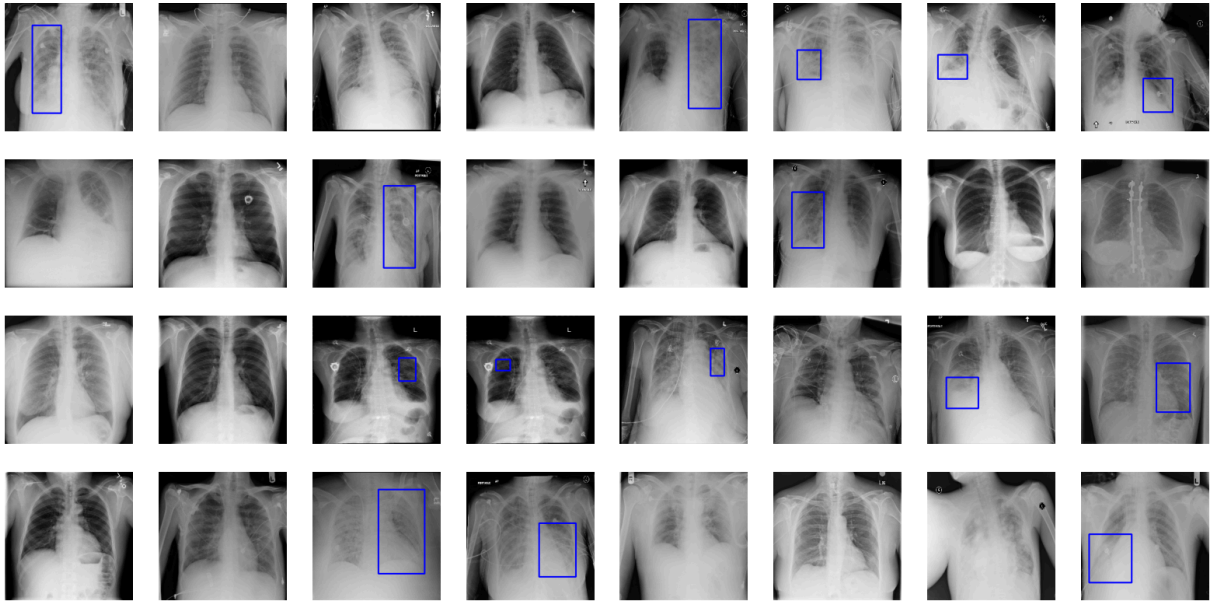
```
In [14]: plt.figure(figsize=(20, 10))
for images, labels in train_ds.take(1):
    for i in range(BATCH_SIZE):
        # print(labels.shape)
        ax = plt.subplot(4, BATCH_SIZE//4, i + 1)
        label = labels[0][i]
        box = (labels[1][i] * input_size)
        box = tf.cast(box, tf.int32)

        image = images[i].numpy().astype("float") * 255.0
        image = image.astype(np.uint8)
        image_color = cv.cvtColor(image, cv.COLOR_GRAY2RGB)

        color = (0, 0, 255)
        if label[0] > 0:
            color = (0, 255, 0)

        cv.rectangle(image_color, box.numpy(), color, 2)
```

```
plt.imshow(image_color)
plt.axis("off")
```



Building model architecture to perform classification (detecting pneumonia) and bounding box regression (localizing pneumonia in X-ray images).

```
In [15]: DROPOUT_FACTOR = 0.5

def build_feature_extractor(inputs):

    x = tf.keras.layers.Conv2D(16, kernel_size=3, activation='relu', input_shape=(i
    x = tf.keras.layers.AveragePooling2D(2,2)(x)

    x = tf.keras.layers.Conv2D(32, kernel_size=3, activation = 'relu')(x)
    x = tf.keras.layers.AveragePooling2D(2,2)(x)

    x = tf.keras.layers.Conv2D(64, kernel_size=3, activation = 'relu')(x)
    x = tf.keras.layers.Dropout(DROPOUT_FACTOR)(x)
    x = tf.keras.layers.AveragePooling2D(2,2)(x)

    return x

def build_model_adaptor(inputs):
    x = tf.keras.layers.Flatten()(inputs)
    x = tf.keras.layers.Dense(64, activation='relu')(x)
    return x

def build_classifier_head(inputs):
    return tf.keras.layers.Dense(CLASSES, activation='softmax', name = 'classifier_

def build_regressor_head(inputs):
    return tf.keras.layers.Dense(units = 4, name = 'regressor_head')(inputs)

def build_model(inputs):

    feature_extractor = build_feature_extractor(inputs)
```

```

model_adaptor = build_model_adaptor(feature_extractor)

classification_head = build_classifier_head(model_adaptor)

regressor_head = build_regressor_head(model_adaptor)

model = tf.keras.Model(inputs = inputs, outputs = [classification_head, regressor_head])

model.compile(optimizer=tf.keras.optimizers.Adam(),
              loss = {'classifier_head' : 'categorical_crossentropy', 'regressor_head' : 'mse'},
              metrics = {'classifier_head' : 'accuracy', 'regressor_head' : 'mse' })

return model

```

Initializing the model with the standardized image dimensions and displaying the model structure.

```

In [16]: model = build_model(tf.keras.layers.Input(shape=(input_size, input_size, 1)))

model.summary()

```

C:\Users\Sanan\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11\_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python311\site-packages\keras\src\layers\convolutional  
base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to  
a layer. When using Sequential models, prefer using an `Input(shape)` object as the  
first layer in the model instead.

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

**Model: "functional"**



```
validation_data=validation_ds, validation_steps=1,  
epochs=100)
```



Epoch 1/100

**187/187** ————— **46s** 216ms/step - classifier\_head\_accuracy: 0.6468 - classifier\_head\_loss: 0.6232 - loss: 0.7324 - regressor\_head\_loss: 0.1093 - regressor\_head\_mse: 0.1093 - val\_classifier\_head\_accuracy: 0.7257 - val\_classifier\_head\_loss: 0.6091 - val\_loss: 0.6427 - val\_regressor\_head\_loss: 0.0336 - val\_regressor\_head\_mse: 0.0336

Epoch 2/100

**187/187** ————— **39s** 210ms/step - classifier\_head\_accuracy: 0.7603 - classifier\_head\_loss: 0.5045 - loss: 0.5396 - regressor\_head\_loss: 0.0351 - regressor\_head\_mse: 0.0351 - val\_classifier\_head\_accuracy: 0.7486 - val\_classifier\_head\_loss: 0.5078 - val\_loss: 0.5263 - val\_regressor\_head\_loss: 0.0185 - val\_regressor\_head\_mse: 0.0185

Epoch 3/100

**187/187** ————— **39s** 211ms/step - classifier\_head\_accuracy: 0.7682 - classifier\_head\_loss: 0.4903 - loss: 0.5217 - regressor\_head\_loss: 0.0314 - regressor\_head\_mse: 0.0314 - val\_classifier\_head\_accuracy: 0.6629 - val\_classifier\_head\_loss: 0.6634 - val\_loss: 0.6998 - val\_regressor\_head\_loss: 0.0364 - val\_regressor\_head\_mse: 0.0364

Epoch 4/100

**187/187** ————— **39s** 208ms/step - classifier\_head\_accuracy: 0.7798 - classifier\_head\_loss: 0.4686 - loss: 0.4985 - regressor\_head\_loss: 0.0300 - regressor\_head\_mse: 0.0300 - val\_classifier\_head\_accuracy: 0.6514 - val\_classifier\_head\_loss: 0.6827 - val\_loss: 0.7124 - val\_regressor\_head\_loss: 0.0297 - val\_regressor\_head\_mse: 0.0297

Epoch 5/100

**187/187** ————— **39s** 210ms/step - classifier\_head\_accuracy: 0.7776 - classifier\_head\_loss: 0.4695 - loss: 0.4998 - regressor\_head\_loss: 0.0303 - regressor\_head\_mse: 0.0303 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 0.5079 - val\_loss: 0.5259 - val\_regressor\_head\_loss: 0.0180 - val\_regressor\_head\_mse: 0.0180

Epoch 6/100

**187/187** ————— **39s** 209ms/step - classifier\_head\_accuracy: 0.7875 - classifier\_head\_loss: 0.4539 - loss: 0.4823 - regressor\_head\_loss: 0.0284 - regressor\_head\_mse: 0.0284 - val\_classifier\_head\_accuracy: 0.8114 - val\_classifier\_head\_loss: 0.4271 - val\_loss: 0.4441 - val\_regressor\_head\_loss: 0.0171 - val\_regressor\_head\_mse: 0.0171

Epoch 7/100

**187/187** ————— **40s** 211ms/step - classifier\_head\_accuracy: 0.7997 - classifier\_head\_loss: 0.4311 - loss: 0.4581 - regressor\_head\_loss: 0.0270 - regressor\_head\_mse: 0.0270 - val\_classifier\_head\_accuracy: 0.8000 - val\_classifier\_head\_loss: 0.4704 - val\_loss: 0.4950 - val\_regressor\_head\_loss: 0.0246 - val\_regressor\_head\_mse: 0.0246

Epoch 8/100

**187/187** ————— **40s** 215ms/step - classifier\_head\_accuracy: 0.8196 - classifier\_head\_loss: 0.4092 - loss: 0.4364 - regressor\_head\_loss: 0.0272 - regressor\_head\_mse: 0.0272 - val\_classifier\_head\_accuracy: 0.7600 - val\_classifier\_head\_loss: 0.4895 - val\_loss: 0.5075 - val\_regressor\_head\_loss: 0.0180 - val\_regressor\_head\_mse: 0.0180

Epoch 9/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.8323 - classifier\_head\_loss: 0.3694 - loss: 0.3955 - regressor\_head\_loss: 0.0261 - regressor\_head\_mse: 0.0261 - val\_classifier\_head\_accuracy: 0.7600 - val\_classifier\_head\_loss: 0.5952 - val\_loss: 0.6190 - val\_regressor\_head\_loss: 0.0238 - val\_regressor\_head\_mse: 0.0238

Epoch 10/100

**187/187** ————— **39s** 211ms/step - classifier\_head\_accuracy: 0.8463 - cla

ssifier\_head\_loss: 0.3521 - loss: 0.3771 - regressor\_head\_loss: 0.0250 - regressor\_head\_mse: 0.0250 - val\_classifier\_head\_accuracy: 0.7371 - val\_classifier\_head\_loss: 0.6611 - val\_loss: 0.6849 - val\_regressor\_head\_loss: 0.0238 - val\_regressor\_head\_mse: 0.0238

Epoch 11/100

**187/187** ————— **42s** 224ms/step - classifier\_head\_accuracy: 0.8576 - classifier\_head\_loss: 0.3304 - loss: 0.3536 - regressor\_head\_loss: 0.0232 - regressor\_head\_mse: 0.0232 - val\_classifier\_head\_accuracy: 0.7200 - val\_classifier\_head\_loss: 0.6706 - val\_loss: 0.6952 - val\_regressor\_head\_loss: 0.0246 - val\_regressor\_head\_mse: 0.0246

Epoch 12/100

**187/187** ————— **43s** 230ms/step - classifier\_head\_accuracy: 0.8727 - classifier\_head\_loss: 0.2956 - loss: 0.3175 - regressor\_head\_loss: 0.0219 - regressor\_head\_mse: 0.0219 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 0.5039 - val\_loss: 0.5236 - val\_regressor\_head\_loss: 0.0196 - val\_regressor\_head\_mse: 0.0196

Epoch 13/100

**187/187** ————— **39s** 210ms/step - classifier\_head\_accuracy: 0.8896 - classifier\_head\_loss: 0.2717 - loss: 0.2916 - regressor\_head\_loss: 0.0199 - regressor\_head\_mse: 0.0199 - val\_classifier\_head\_accuracy: 0.8286 - val\_classifier\_head\_loss: 0.4894 - val\_loss: 0.5090 - val\_regressor\_head\_loss: 0.0196 - val\_regressor\_head\_mse: 0.0196

Epoch 14/100

**187/187** ————— **39s** 210ms/step - classifier\_head\_accuracy: 0.8915 - classifier\_head\_loss: 0.2516 - loss: 0.2726 - regressor\_head\_loss: 0.0211 - regressor\_head\_mse: 0.0211 - val\_classifier\_head\_accuracy: 0.8171 - val\_classifier\_head\_loss: 0.4354 - val\_loss: 0.4537 - val\_regressor\_head\_loss: 0.0183 - val\_regressor\_head\_mse: 0.0183

Epoch 15/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9046 - classifier\_head\_loss: 0.2220 - loss: 0.2432 - regressor\_head\_loss: 0.0212 - regressor\_head\_mse: 0.0212 - val\_classifier\_head\_accuracy: 0.7486 - val\_classifier\_head\_loss: 0.6907 - val\_loss: 0.7143 - val\_regressor\_head\_loss: 0.0236 - val\_regressor\_head\_mse: 0.0236

Epoch 16/100

**187/187** ————— **39s** 210ms/step - classifier\_head\_accuracy: 0.9155 - classifier\_head\_loss: 0.2036 - loss: 0.2242 - regressor\_head\_loss: 0.0206 - regressor\_head\_mse: 0.0206 - val\_classifier\_head\_accuracy: 0.7543 - val\_classifier\_head\_loss: 0.6855 - val\_loss: 0.7110 - val\_regressor\_head\_loss: 0.0255 - val\_regressor\_head\_mse: 0.0255

Epoch 17/100

**187/187** ————— **39s** 211ms/step - classifier\_head\_accuracy: 0.9185 - classifier\_head\_loss: 0.1870 - loss: 0.2075 - regressor\_head\_loss: 0.0205 - regressor\_head\_mse: 0.0205 - val\_classifier\_head\_accuracy: 0.6514 - val\_classifier\_head\_loss: 1.0186 - val\_loss: 1.0488 - val\_regressor\_head\_loss: 0.0303 - val\_regressor\_head\_mse: 0.0303

Epoch 18/100

**187/187** ————— **40s** 211ms/step - classifier\_head\_accuracy: 0.9286 - classifier\_head\_loss: 0.1714 - loss: 0.1916 - regressor\_head\_loss: 0.0202 - regressor\_head\_mse: 0.0202 - val\_classifier\_head\_accuracy: 0.6971 - val\_classifier\_head\_loss: 0.8996 - val\_loss: 0.9282 - val\_regressor\_head\_loss: 0.0286 - val\_regressor\_head\_mse: 0.0286

Epoch 19/100

**187/187** ————— **39s** 210ms/step - classifier\_head\_accuracy: 0.9359 - classifier\_head\_loss: 0.1538 - loss: 0.1736 - regressor\_head\_loss: 0.0198 - regressor\_head\_mse: 0.0198 - val\_classifier\_head\_accuracy: 0.8114 - val\_classifier\_head\_loss:

0.6495 - val\_loss: 0.6691 - val\_regressor\_head\_loss: 0.0196 - val\_regressor\_head\_mse: 0.0196  
Epoch 20/100  
**187/187** ————— 39s 211ms/step - classifier\_head\_accuracy: 0.9410 - classifier\_head\_loss: 0.1401 - loss: 0.1590 - regressor\_head\_loss: 0.0189 - regressor\_head\_mse: 0.0189 - val\_classifier\_head\_accuracy: 0.7657 - val\_classifier\_head\_loss: 0.7286 - val\_loss: 0.7482 - val\_regressor\_head\_loss: 0.0196 - val\_regressor\_head\_mse: 0.0196  
Epoch 21/100  
**187/187** ————— 40s 212ms/step - classifier\_head\_accuracy: 0.9459 - classifier\_head\_loss: 0.1261 - loss: 0.1456 - regressor\_head\_loss: 0.0195 - regressor\_head\_mse: 0.0195 - val\_classifier\_head\_accuracy: 0.7600 - val\_classifier\_head\_loss: 0.9365 - val\_loss: 0.9600 - val\_regressor\_head\_loss: 0.0235 - val\_regressor\_head\_mse: 0.0235  
Epoch 22/100  
**187/187** ————— 40s 212ms/step - classifier\_head\_accuracy: 0.9503 - classifier\_head\_loss: 0.1123 - loss: 0.1315 - regressor\_head\_loss: 0.0192 - regressor\_head\_mse: 0.0192 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.0586 - val\_loss: 1.0834 - val\_regressor\_head\_loss: 0.0248 - val\_regressor\_head\_mse: 0.0248  
Epoch 23/100  
**187/187** ————— 40s 213ms/step - classifier\_head\_accuracy: 0.9627 - classifier\_head\_loss: 0.0936 - loss: 0.1125 - regressor\_head\_loss: 0.0189 - regressor\_head\_mse: 0.0189 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.1600 - val\_loss: 1.1850 - val\_regressor\_head\_loss: 0.0250 - val\_regressor\_head\_mse: 0.0250  
Epoch 24/100  
**187/187** ————— 40s 214ms/step - classifier\_head\_accuracy: 0.9596 - classifier\_head\_loss: 0.0967 - loss: 0.1159 - regressor\_head\_loss: 0.0192 - regressor\_head\_mse: 0.0192 - val\_classifier\_head\_accuracy: 0.7200 - val\_classifier\_head\_loss: 1.0390 - val\_loss: 1.0637 - val\_regressor\_head\_loss: 0.0246 - val\_regressor\_head\_mse: 0.0246  
Epoch 25/100  
**187/187** ————— 39s 211ms/step - classifier\_head\_accuracy: 0.9560 - classifier\_head\_loss: 0.0985 - loss: 0.1170 - regressor\_head\_loss: 0.0185 - regressor\_head\_mse: 0.0185 - val\_classifier\_head\_accuracy: 0.7657 - val\_classifier\_head\_loss: 0.9505 - val\_loss: 0.9728 - val\_regressor\_head\_loss: 0.0223 - val\_regressor\_head\_mse: 0.0223  
Epoch 26/100  
**187/187** ————— 40s 211ms/step - classifier\_head\_accuracy: 0.9680 - classifier\_head\_loss: 0.0797 - loss: 0.0974 - regressor\_head\_loss: 0.0177 - regressor\_head\_mse: 0.0177 - val\_classifier\_head\_accuracy: 0.8229 - val\_classifier\_head\_loss: 0.6768 - val\_loss: 0.6943 - val\_regressor\_head\_loss: 0.0175 - val\_regressor\_head\_mse: 0.0175  
Epoch 27/100  
**187/187** ————— 39s 211ms/step - classifier\_head\_accuracy: 0.9669 - classifier\_head\_loss: 0.0795 - loss: 0.0978 - regressor\_head\_loss: 0.0183 - regressor\_head\_mse: 0.0183 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss: 1.0299 - val\_loss: 1.0523 - val\_regressor\_head\_loss: 0.0224 - val\_regressor\_head\_mse: 0.0224  
Epoch 28/100  
**187/187** ————— 39s 211ms/step - classifier\_head\_accuracy: 0.9759 - classifier\_head\_loss: 0.0623 - loss: 0.0804 - regressor\_head\_loss: 0.0181 - regressor\_head\_mse: 0.0181 - val\_classifier\_head\_accuracy: 0.7143 - val\_classifier\_head\_loss: 1.5558 - val\_loss: 1.5805 - val\_regressor\_head\_loss: 0.0247 - val\_regressor\_head\_mse: 0.0247

Epoch 29/100

**187/187** ————— **40s** 211ms/step - classifier\_head\_accuracy: 0.9766 - classifier\_head\_loss: 0.0607 - loss: 0.0799 - regressor\_head\_loss: 0.0192 - regressor\_head\_mse: 0.0192 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.3288 - val\_loss: 1.3517 - val\_regressor\_head\_loss: 0.0229 - val\_regressor\_head\_mse: 0.0229

Epoch 30/100

**187/187** ————— **40s** 213ms/step - classifier\_head\_accuracy: 0.9741 - classifier\_head\_loss: 0.0652 - loss: 0.0839 - regressor\_head\_loss: 0.0187 - regressor\_head\_mse: 0.0187 - val\_classifier\_head\_accuracy: 0.7943 - val\_classifier\_head\_loss: 1.0911 - val\_loss: 1.1119 - val\_regressor\_head\_loss: 0.0208 - val\_regressor\_head\_mse: 0.0208

Epoch 31/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9631 - classifier\_head\_loss: 0.0898 - loss: 0.1093 - regressor\_head\_loss: 0.0195 - regressor\_head\_mse: 0.0195 - val\_classifier\_head\_accuracy: 0.7086 - val\_classifier\_head\_loss: 1.7324 - val\_loss: 1.7611 - val\_regressor\_head\_loss: 0.0286 - val\_regressor\_head\_mse: 0.0286

Epoch 32/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9743 - classifier\_head\_loss: 0.0616 - loss: 0.0796 - regressor\_head\_loss: 0.0179 - regressor\_head\_mse: 0.0179 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss: 1.1491 - val\_loss: 1.1705 - val\_regressor\_head\_loss: 0.0215 - val\_regressor\_head\_mse: 0.0215

Epoch 33/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9807 - classifier\_head\_loss: 0.0486 - loss: 0.0661 - regressor\_head\_loss: 0.0175 - regressor\_head\_mse: 0.0175 - val\_classifier\_head\_accuracy: 0.7886 - val\_classifier\_head\_loss: 1.1015 - val\_loss: 1.1209 - val\_regressor\_head\_loss: 0.0194 - val\_regressor\_head\_mse: 0.0194

Epoch 34/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9819 - classifier\_head\_loss: 0.0477 - loss: 0.0654 - regressor\_head\_loss: 0.0178 - regressor\_head\_mse: 0.0178 - val\_classifier\_head\_accuracy: 0.7943 - val\_classifier\_head\_loss: 1.3595 - val\_loss: 1.3809 - val\_regressor\_head\_loss: 0.0214 - val\_regressor\_head\_mse: 0.0214

Epoch 35/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9810 - classifier\_head\_loss: 0.0543 - loss: 0.0727 - regressor\_head\_loss: 0.0183 - regressor\_head\_mse: 0.0183 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.5935 - val\_loss: 1.6174 - val\_regressor\_head\_loss: 0.0239 - val\_regressor\_head\_mse: 0.0239

Epoch 36/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9782 - classifier\_head\_loss: 0.0566 - loss: 0.0748 - regressor\_head\_loss: 0.0182 - regressor\_head\_mse: 0.0182 - val\_classifier\_head\_accuracy: 0.7657 - val\_classifier\_head\_loss: 1.4251 - val\_loss: 1.4471 - val\_regressor\_head\_loss: 0.0220 - val\_regressor\_head\_mse: 0.0220

Epoch 37/100

**187/187** ————— **39s** 211ms/step - classifier\_head\_accuracy: 0.9774 - classifier\_head\_loss: 0.0500 - loss: 0.0680 - regressor\_head\_loss: 0.0180 - regressor\_head\_mse: 0.0180 - val\_classifier\_head\_accuracy: 0.7543 - val\_classifier\_head\_loss: 1.5222 - val\_loss: 1.5457 - val\_regressor\_head\_loss: 0.0235 - val\_regressor\_head\_mse: 0.0235

Epoch 38/100

**187/187** ————— **40s** 214ms/step - classifier\_head\_accuracy: 0.9872 - cla

ssifier\_head\_loss: 0.0369 - loss: 0.0548 - regressor\_head\_loss: 0.0179 - regressor\_head\_mse: 0.0179 - val\_classifier\_head\_accuracy: 0.7543 - val\_classifier\_head\_loss: 1.4985 - val\_loss: 1.5231 - val\_regressor\_head\_loss: 0.0246 - val\_regressor\_head\_mse: 0.0246

Epoch 39/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9837 - classifier\_head\_loss: 0.0395 - loss: 0.0563 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.7086 - val\_classifier\_head\_loss: 1.6551 - val\_loss: 1.6777 - val\_regressor\_head\_loss: 0.0227 - val\_regressor\_head\_mse: 0.0227

Epoch 40/100

**187/187** ————— **40s** 214ms/step - classifier\_head\_accuracy: 0.9881 - classifier\_head\_loss: 0.0318 - loss: 0.0485 - regressor\_head\_loss: 0.0167 - regressor\_head\_mse: 0.0167 - val\_classifier\_head\_accuracy: 0.6914 - val\_classifier\_head\_loss: 1.9072 - val\_loss: 1.9326 - val\_regressor\_head\_loss: 0.0255 - val\_regressor\_head\_mse: 0.0255

Epoch 41/100

**187/187** ————— **40s** 215ms/step - classifier\_head\_accuracy: 0.9887 - classifier\_head\_loss: 0.0313 - loss: 0.0483 - regressor\_head\_loss: 0.0170 - regressor\_head\_mse: 0.0170 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.5979 - val\_loss: 1.6225 - val\_regressor\_head\_loss: 0.0246 - val\_regressor\_head\_mse: 0.0246

Epoch 42/100

**187/187** ————— **41s** 217ms/step - classifier\_head\_accuracy: 0.9848 - classifier\_head\_loss: 0.0409 - loss: 0.0588 - regressor\_head\_loss: 0.0179 - regressor\_head\_mse: 0.0179 - val\_classifier\_head\_accuracy: 0.7657 - val\_classifier\_head\_loss: 1.7820 - val\_loss: 1.8067 - val\_regressor\_head\_loss: 0.0247 - val\_regressor\_head\_mse: 0.0247

Epoch 43/100

**187/187** ————— **40s** 215ms/step - classifier\_head\_accuracy: 0.9843 - classifier\_head\_loss: 0.0400 - loss: 0.0577 - regressor\_head\_loss: 0.0177 - regressor\_head\_mse: 0.0177 - val\_classifier\_head\_accuracy: 0.7771 - val\_classifier\_head\_loss: 1.5309 - val\_loss: 1.5536 - val\_regressor\_head\_loss: 0.0227 - val\_regressor\_head\_mse: 0.0227

Epoch 44/100

**187/187** ————— **40s** 213ms/step - classifier\_head\_accuracy: 0.9843 - classifier\_head\_loss: 0.0430 - loss: 0.0609 - regressor\_head\_loss: 0.0179 - regressor\_head\_mse: 0.0179 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.7636 - val\_loss: 1.7875 - val\_regressor\_head\_loss: 0.0239 - val\_regressor\_head\_mse: 0.0239

Epoch 45/100

**187/187** ————— **40s** 213ms/step - classifier\_head\_accuracy: 0.9856 - classifier\_head\_loss: 0.0342 - loss: 0.0512 - regressor\_head\_loss: 0.0170 - regressor\_head\_mse: 0.0170 - val\_classifier\_head\_accuracy: 0.7371 - val\_classifier\_head\_loss: 1.5928 - val\_loss: 1.6188 - val\_regressor\_head\_loss: 0.0260 - val\_regressor\_head\_mse: 0.0260

Epoch 46/100

**187/187** ————— **40s** 213ms/step - classifier\_head\_accuracy: 0.9832 - classifier\_head\_loss: 0.0459 - loss: 0.0623 - regressor\_head\_loss: 0.0164 - regressor\_head\_mse: 0.0164 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss: 1.2363 - val\_loss: 1.2575 - val\_regressor\_head\_loss: 0.0212 - val\_regressor\_head\_mse: 0.0212

Epoch 47/100

**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9807 - classifier\_head\_loss: 0.0500 - loss: 0.0670 - regressor\_head\_loss: 0.0171 - regressor\_head\_mse: 0.0171 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss:

1.5888 - val\_loss: 1.6130 - val\_regressor\_head\_loss: 0.0243 - val\_regressor\_head\_mse: 0.0243  
Epoch 48/100  
**187/187** ————— **40s** 214ms/step - classifier\_head\_accuracy: 0.9861 - classifier\_head\_loss: 0.0384 - loss: 0.0559 - regressor\_head\_loss: 0.0175 - regressor\_head\_mse: 0.0175 - val\_classifier\_head\_accuracy: 0.6914 - val\_classifier\_head\_loss: 2.4200 - val\_loss: 2.4500 - val\_regressor\_head\_loss: 0.0300 - val\_regressor\_head\_mse: 0.0300  
Epoch 49/100  
**187/187** ————— **41s** 217ms/step - classifier\_head\_accuracy: 0.9878 - classifier\_head\_loss: 0.0309 - loss: 0.0482 - regressor\_head\_loss: 0.0173 - regressor\_head\_mse: 0.0173 - val\_classifier\_head\_accuracy: 0.7600 - val\_classifier\_head\_loss: 1.4386 - val\_loss: 1.4632 - val\_regressor\_head\_loss: 0.0246 - val\_regressor\_head\_mse: 0.0246  
Epoch 50/100  
**187/187** ————— **40s** 212ms/step - classifier\_head\_accuracy: 0.9832 - classifier\_head\_loss: 0.0470 - loss: 0.0647 - regressor\_head\_loss: 0.0177 - regressor\_head\_mse: 0.0177 - val\_classifier\_head\_accuracy: 0.7314 - val\_classifier\_head\_loss: 1.7721 - val\_loss: 1.7955 - val\_regressor\_head\_loss: 0.0233 - val\_regressor\_head\_mse: 0.0233  
Epoch 51/100  
**187/187** ————— **40s** 213ms/step - classifier\_head\_accuracy: 0.9903 - classifier\_head\_loss: 0.0230 - loss: 0.0392 - regressor\_head\_loss: 0.0162 - regressor\_head\_mse: 0.0162 - val\_classifier\_head\_accuracy: 0.6629 - val\_classifier\_head\_loss: 2.1473 - val\_loss: 2.1723 - val\_regressor\_head\_loss: 0.0250 - val\_regressor\_head\_mse: 0.0250  
Epoch 52/100  
**187/187** ————— **40s** 213ms/step - classifier\_head\_accuracy: 0.9843 - classifier\_head\_loss: 0.0364 - loss: 0.0528 - regressor\_head\_loss: 0.0164 - regressor\_head\_mse: 0.0164 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.6558 - val\_loss: 1.6782 - val\_regressor\_head\_loss: 0.0225 - val\_regressor\_head\_mse: 0.0225  
Epoch 53/100  
**187/187** ————— **40s** 214ms/step - classifier\_head\_accuracy: 0.9881 - classifier\_head\_loss: 0.0363 - loss: 0.0524 - regressor\_head\_loss: 0.0161 - regressor\_head\_mse: 0.0161 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss: 1.4727 - val\_loss: 1.4945 - val\_regressor\_head\_loss: 0.0218 - val\_regressor\_head\_mse: 0.0218  
Epoch 54/100  
**187/187** ————— **40s** 215ms/step - classifier\_head\_accuracy: 0.9892 - classifier\_head\_loss: 0.0312 - loss: 0.0480 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.9103 - val\_loss: 1.9365 - val\_regressor\_head\_loss: 0.0262 - val\_regressor\_head\_mse: 0.0262  
Epoch 55/100  
**187/187** ————— **40s** 213ms/step - classifier\_head\_accuracy: 0.9899 - classifier\_head\_loss: 0.0235 - loss: 0.0403 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.7086 - val\_classifier\_head\_loss: 2.0845 - val\_loss: 2.1087 - val\_regressor\_head\_loss: 0.0242 - val\_regressor\_head\_mse: 0.0242  
Epoch 56/100  
**187/187** ————— **41s** 217ms/step - classifier\_head\_accuracy: 0.9876 - classifier\_head\_loss: 0.0321 - loss: 0.0490 - regressor\_head\_loss: 0.0169 - regressor\_head\_mse: 0.0169 - val\_classifier\_head\_accuracy: 0.7657 - val\_classifier\_head\_loss: 1.2885 - val\_loss: 1.3101 - val\_regressor\_head\_loss: 0.0216 - val\_regressor\_head\_mse: 0.0216



Epoch 57/100

**187/187** ————— **41s** 219ms/step - classifier\_head\_accuracy: 0.9874 - classifier\_head\_loss: 0.0369 - loss: 0.0536 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.8000 - val\_classifier\_head\_loss: 1.4535 - val\_loss: 1.4765 - val\_regressor\_head\_loss: 0.0230 - val\_regressor\_head\_mse: 0.0230

Epoch 58/100

**187/187** ————— **41s** 219ms/step - classifier\_head\_accuracy: 0.9857 - classifier\_head\_loss: 0.0352 - loss: 0.0522 - regressor\_head\_loss: 0.0169 - regressor\_head\_mse: 0.0169 - val\_classifier\_head\_accuracy: 0.7543 - val\_classifier\_head\_loss: 1.8257 - val\_loss: 1.8501 - val\_regressor\_head\_loss: 0.0243 - val\_regressor\_head\_mse: 0.0243

Epoch 59/100

**187/187** ————— **41s** 218ms/step - classifier\_head\_accuracy: 0.9882 - classifier\_head\_loss: 0.0289 - loss: 0.0442 - regressor\_head\_loss: 0.0153 - regressor\_head\_mse: 0.0153 - val\_classifier\_head\_accuracy: 0.7943 - val\_classifier\_head\_loss: 1.5359 - val\_loss: 1.5569 - val\_regressor\_head\_loss: 0.0210 - val\_regressor\_head\_mse: 0.0210

Epoch 60/100

**187/187** ————— **41s** 219ms/step - classifier\_head\_accuracy: 0.9870 - classifier\_head\_loss: 0.0336 - loss: 0.0500 - regressor\_head\_loss: 0.0164 - regressor\_head\_mse: 0.0164 - val\_classifier\_head\_accuracy: 0.7371 - val\_classifier\_head\_loss: 2.0457 - val\_loss: 2.0709 - val\_regressor\_head\_loss: 0.0251 - val\_regressor\_head\_mse: 0.0251

Epoch 61/100

**187/187** ————— **41s** 220ms/step - classifier\_head\_accuracy: 0.9878 - classifier\_head\_loss: 0.0307 - loss: 0.0475 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.7886 - val\_classifier\_head\_loss: 1.6897 - val\_loss: 1.7127 - val\_regressor\_head\_loss: 0.0230 - val\_regressor\_head\_mse: 0.0230

Epoch 62/100

**187/187** ————— **41s** 220ms/step - classifier\_head\_accuracy: 0.9925 - classifier\_head\_loss: 0.0196 - loss: 0.0364 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.8872 - val\_loss: 1.9099 - val\_regressor\_head\_loss: 0.0227 - val\_regressor\_head\_mse: 0.0227

Epoch 63/100

**187/187** ————— **41s** 220ms/step - classifier\_head\_accuracy: 0.9900 - classifier\_head\_loss: 0.0268 - loss: 0.0432 - regressor\_head\_loss: 0.0164 - regressor\_head\_mse: 0.0164 - val\_classifier\_head\_accuracy: 0.7257 - val\_classifier\_head\_loss: 2.0900 - val\_loss: 2.1133 - val\_regressor\_head\_loss: 0.0233 - val\_regressor\_head\_mse: 0.0233

Epoch 64/100

**187/187** ————— **41s** 221ms/step - classifier\_head\_accuracy: 0.9883 - classifier\_head\_loss: 0.0281 - loss: 0.0447 - regressor\_head\_loss: 0.0166 - regressor\_head\_mse: 0.0166 - val\_classifier\_head\_accuracy: 0.7886 - val\_classifier\_head\_loss: 1.5127 - val\_loss: 1.5365 - val\_regressor\_head\_loss: 0.0238 - val\_regressor\_head\_mse: 0.0238

Epoch 65/100

**187/187** ————— **41s** 222ms/step - classifier\_head\_accuracy: 0.9866 - classifier\_head\_loss: 0.0361 - loss: 0.0530 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.7600 - val\_classifier\_head\_loss: 1.6834 - val\_loss: 1.7091 - val\_regressor\_head\_loss: 0.0257 - val\_regressor\_head\_mse: 0.0257

Epoch 66/100

**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9841 - cla

ssifier\_head\_loss: 0.0440 - loss: 0.0602 - regressor\_head\_loss: 0.0162 - regressor\_head\_mse: 0.0162 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.6038 - val\_loss: 1.6273 - val\_regressor\_head\_loss: 0.0235 - val\_regressor\_head\_mse: 0.0235

Epoch 67/100

**187/187** ————— **41s** 221ms/step - classifier\_head\_accuracy: 0.9910 - classifier\_head\_loss: 0.0218 - loss: 0.0375 - regressor\_head\_loss: 0.0157 - regressor\_head\_mse: 0.0157 - val\_classifier\_head\_accuracy: 0.7943 - val\_classifier\_head\_loss: 1.7083 - val\_loss: 1.7297 - val\_regressor\_head\_loss: 0.0213 - val\_regressor\_head\_mse: 0.0213

Epoch 68/100

**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9914 - classifier\_head\_loss: 0.0216 - loss: 0.0374 - regressor\_head\_loss: 0.0158 - regressor\_head\_mse: 0.0158 - val\_classifier\_head\_accuracy: 0.8000 - val\_classifier\_head\_loss: 1.1383 - val\_loss: 1.1571 - val\_regressor\_head\_loss: 0.0188 - val\_regressor\_head\_mse: 0.0188

Epoch 69/100

**187/187** ————— **41s** 221ms/step - classifier\_head\_accuracy: 0.9857 - classifier\_head\_loss: 0.0404 - loss: 0.0572 - regressor\_head\_loss: 0.0168 - regressor\_head\_mse: 0.0168 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.9426 - val\_loss: 1.9648 - val\_regressor\_head\_loss: 0.0222 - val\_regressor\_head\_mse: 0.0222

Epoch 70/100

**187/187** ————— **41s** 219ms/step - classifier\_head\_accuracy: 0.9905 - classifier\_head\_loss: 0.0259 - loss: 0.0424 - regressor\_head\_loss: 0.0165 - regressor\_head\_mse: 0.0165 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.7696 - val\_loss: 1.7928 - val\_regressor\_head\_loss: 0.0233 - val\_regressor\_head\_mse: 0.0233

Epoch 71/100

**187/187** ————— **41s** 222ms/step - classifier\_head\_accuracy: 0.9931 - classifier\_head\_loss: 0.0162 - loss: 0.0320 - regressor\_head\_loss: 0.0158 - regressor\_head\_mse: 0.0158 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.7214 - val\_loss: 1.7413 - val\_regressor\_head\_loss: 0.0199 - val\_regressor\_head\_mse: 0.0199

Epoch 72/100

**187/187** ————— **41s** 221ms/step - classifier\_head\_accuracy: 0.9896 - classifier\_head\_loss: 0.0256 - loss: 0.0410 - regressor\_head\_loss: 0.0153 - regressor\_head\_mse: 0.0153 - val\_classifier\_head\_accuracy: 0.8057 - val\_classifier\_head\_loss: 1.5728 - val\_loss: 1.5931 - val\_regressor\_head\_loss: 0.0204 - val\_regressor\_head\_mse: 0.0204

Epoch 73/100

**187/187** ————— **41s** 221ms/step - classifier\_head\_accuracy: 0.9892 - classifier\_head\_loss: 0.0294 - loss: 0.0450 - regressor\_head\_loss: 0.0156 - regressor\_head\_mse: 0.0156 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.4767 - val\_loss: 1.4969 - val\_regressor\_head\_loss: 0.0201 - val\_regressor\_head\_mse: 0.0201

Epoch 74/100

**187/187** ————— **42s** 222ms/step - classifier\_head\_accuracy: 0.9909 - classifier\_head\_loss: 0.0263 - loss: 0.0425 - regressor\_head\_loss: 0.0162 - regressor\_head\_mse: 0.0162 - val\_classifier\_head\_accuracy: 0.7771 - val\_classifier\_head\_loss: 1.7613 - val\_loss: 1.7829 - val\_regressor\_head\_loss: 0.0216 - val\_regressor\_head\_mse: 0.0216

Epoch 75/100

**187/187** ————— **41s** 222ms/step - classifier\_head\_accuracy: 0.9907 - classifier\_head\_loss: 0.0237 - loss: 0.0400 - regressor\_head\_loss: 0.0163 - regressor\_head\_mse: 0.0163 - val\_classifier\_head\_accuracy: 0.8057 - val\_classifier\_head\_loss:



1.5534 - val\_loss: 1.5750 - val\_regressor\_head\_loss: 0.0217 - val\_regressor\_head\_mse: 0.0217  
Epoch 76/100  
**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9913 - classifier\_head\_loss: 0.0210 - loss: 0.0371 - regressor\_head\_loss: 0.0161 - regressor\_head\_mse: 0.0161 - val\_classifier\_head\_accuracy: 0.7771 - val\_classifier\_head\_loss: 1.8577 - val\_loss: 1.8808 - val\_regressor\_head\_loss: 0.0231 - val\_regressor\_head\_mse: 0.0231  
Epoch 77/100  
**187/187** ————— **42s** 224ms/step - classifier\_head\_accuracy: 0.9902 - classifier\_head\_loss: 0.0259 - loss: 0.0423 - regressor\_head\_loss: 0.0165 - regressor\_head\_mse: 0.0165 - val\_classifier\_head\_accuracy: 0.6857 - val\_classifier\_head\_loss: 2.7404 - val\_loss: 2.7652 - val\_regressor\_head\_loss: 0.0248 - val\_regressor\_head\_mse: 0.0248  
Epoch 78/100  
**187/187** ————— **42s** 225ms/step - classifier\_head\_accuracy: 0.9881 - classifier\_head\_loss: 0.0227 - loss: 0.0380 - regressor\_head\_loss: 0.0153 - regressor\_head\_mse: 0.0153 - val\_classifier\_head\_accuracy: 0.7257 - val\_classifier\_head\_loss: 2.2765 - val\_loss: 2.3002 - val\_regressor\_head\_loss: 0.0237 - val\_regressor\_head\_mse: 0.0237  
Epoch 79/100  
**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9898 - classifier\_head\_loss: 0.0231 - loss: 0.0378 - regressor\_head\_loss: 0.0146 - regressor\_head\_mse: 0.0146 - val\_classifier\_head\_accuracy: 0.8171 - val\_classifier\_head\_loss: 1.4309 - val\_loss: 1.4506 - val\_regressor\_head\_loss: 0.0198 - val\_regressor\_head\_mse: 0.0198  
Epoch 80/100  
**187/187** ————— **42s** 224ms/step - classifier\_head\_accuracy: 0.9901 - classifier\_head\_loss: 0.0279 - loss: 0.0434 - regressor\_head\_loss: 0.0155 - regressor\_head\_mse: 0.0155 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss: 1.4946 - val\_loss: 1.5150 - val\_regressor\_head\_loss: 0.0204 - val\_regressor\_head\_mse: 0.0204  
Epoch 81/100  
**187/187** ————— **42s** 225ms/step - classifier\_head\_accuracy: 0.9901 - classifier\_head\_loss: 0.0233 - loss: 0.0387 - regressor\_head\_loss: 0.0154 - regressor\_head\_mse: 0.0154 - val\_classifier\_head\_accuracy: 0.7086 - val\_classifier\_head\_loss: 2.0832 - val\_loss: 2.1091 - val\_regressor\_head\_loss: 0.0259 - val\_regressor\_head\_mse: 0.0259  
Epoch 82/100  
**187/187** ————— **42s** 225ms/step - classifier\_head\_accuracy: 0.9916 - classifier\_head\_loss: 0.0202 - loss: 0.0359 - regressor\_head\_loss: 0.0157 - regressor\_head\_mse: 0.0157 - val\_classifier\_head\_accuracy: 0.7771 - val\_classifier\_head\_loss: 1.6226 - val\_loss: 1.6419 - val\_regressor\_head\_loss: 0.0192 - val\_regressor\_head\_mse: 0.0192  
Epoch 83/100  
**187/187** ————— **42s** 225ms/step - classifier\_head\_accuracy: 0.9929 - classifier\_head\_loss: 0.0178 - loss: 0.0331 - regressor\_head\_loss: 0.0154 - regressor\_head\_mse: 0.0154 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.6744 - val\_loss: 1.6945 - val\_regressor\_head\_loss: 0.0201 - val\_regressor\_head\_mse: 0.0201  
Epoch 84/100  
**187/187** ————— **43s** 230ms/step - classifier\_head\_accuracy: 0.9912 - classifier\_head\_loss: 0.0171 - loss: 0.0321 - regressor\_head\_loss: 0.0150 - regressor\_head\_mse: 0.0150 - val\_classifier\_head\_accuracy: 0.7371 - val\_classifier\_head\_loss: 2.3194 - val\_loss: 2.3408 - val\_regressor\_head\_loss: 0.0214 - val\_regressor\_head\_mse: 0.0214

Epoch 85/100

**187/187** ————— **42s** 225ms/step - classifier\_head\_accuracy: 0.9897 - classifier\_head\_loss: 0.0215 - loss: 0.0366 - regressor\_head\_loss: 0.0151 - regressor\_head\_mse: 0.0151 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 1.8041 - val\_loss: 1.8260 - val\_regressor\_head\_loss: 0.0218 - val\_regressor\_head\_mse: 0.0218

Epoch 86/100

**187/187** ————— **41s** 221ms/step - classifier\_head\_accuracy: 0.9903 - classifier\_head\_loss: 0.0201 - loss: 0.0344 - regressor\_head\_loss: 0.0143 - regressor\_head\_mse: 0.0143 - val\_classifier\_head\_accuracy: 0.6971 - val\_classifier\_head\_loss: 2.3628 - val\_loss: 2.3851 - val\_regressor\_head\_loss: 0.0223 - val\_regressor\_head\_mse: 0.0223

Epoch 87/100

**187/187** ————— **42s** 222ms/step - classifier\_head\_accuracy: 0.9883 - classifier\_head\_loss: 0.0404 - loss: 0.0559 - regressor\_head\_loss: 0.0155 - regressor\_head\_mse: 0.0155 - val\_classifier\_head\_accuracy: 0.8171 - val\_classifier\_head\_loss: 1.5490 - val\_loss: 1.5697 - val\_regressor\_head\_loss: 0.0207 - val\_regressor\_head\_mse: 0.0207

Epoch 88/100

**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9914 - classifier\_head\_loss: 0.0225 - loss: 0.0381 - regressor\_head\_loss: 0.0157 - regressor\_head\_mse: 0.0157 - val\_classifier\_head\_accuracy: 0.7943 - val\_classifier\_head\_loss: 1.6218 - val\_loss: 1.6419 - val\_regressor\_head\_loss: 0.0201 - val\_regressor\_head\_mse: 0.0201

Epoch 89/100

**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9943 - classifier\_head\_loss: 0.0140 - loss: 0.0293 - regressor\_head\_loss: 0.0154 - regressor\_head\_mse: 0.0154 - val\_classifier\_head\_accuracy: 0.7714 - val\_classifier\_head\_loss: 1.8369 - val\_loss: 1.8575 - val\_regressor\_head\_loss: 0.0206 - val\_regressor\_head\_mse: 0.0206

Epoch 90/100

**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9872 - classifier\_head\_loss: 0.0284 - loss: 0.0444 - regressor\_head\_loss: 0.0160 - regressor\_head\_mse: 0.0160 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss: 1.9429 - val\_loss: 1.9614 - val\_regressor\_head\_loss: 0.0185 - val\_regressor\_head\_mse: 0.0185

Epoch 91/100

**187/187** ————— **42s** 223ms/step - classifier\_head\_accuracy: 0.9893 - classifier\_head\_loss: 0.0236 - loss: 0.0393 - regressor\_head\_loss: 0.0156 - regressor\_head\_mse: 0.0156 - val\_classifier\_head\_accuracy: 0.7829 - val\_classifier\_head\_loss: 1.7864 - val\_loss: 1.8084 - val\_regressor\_head\_loss: 0.0220 - val\_regressor\_head\_mse: 0.0220

Epoch 92/100

**187/187** ————— **42s** 224ms/step - classifier\_head\_accuracy: 0.9908 - classifier\_head\_loss: 0.0204 - loss: 0.0354 - regressor\_head\_loss: 0.0150 - regressor\_head\_mse: 0.0150 - val\_classifier\_head\_accuracy: 0.8229 - val\_classifier\_head\_loss: 1.4334 - val\_loss: 1.4527 - val\_regressor\_head\_loss: 0.0193 - val\_regressor\_head\_mse: 0.0193

Epoch 93/100

**187/187** ————— **42s** 224ms/step - classifier\_head\_accuracy: 0.9905 - classifier\_head\_loss: 0.0222 - loss: 0.0366 - regressor\_head\_loss: 0.0144 - regressor\_head\_mse: 0.0144 - val\_classifier\_head\_accuracy: 0.8343 - val\_classifier\_head\_loss: 1.1606 - val\_loss: 1.1783 - val\_regressor\_head\_loss: 0.0177 - val\_regressor\_head\_mse: 0.0177

Epoch 94/100

**187/187** ————— **42s** 224ms/step - classifier\_head\_accuracy: 0.9914 - cla

ssifier\_head\_loss: 0.0207 - loss: 0.0358 - regressor\_head\_loss: 0.0151 - regressor\_head\_mse: 0.0151 - val\_classifier\_head\_accuracy: 0.8229 - val\_classifier\_head\_loss: 1.6305 - val\_loss: 1.6514 - val\_regressor\_head\_loss: 0.0209 - val\_regressor\_head\_mse: 0.0209

Epoch 95/100

**187/187** ————— 42s 224ms/step - classifier\_head\_accuracy: 0.9948 - classifier\_head\_loss: 0.0164 - loss: 0.0317 - regressor\_head\_loss: 0.0152 - regressor\_head\_mse: 0.0152 - val\_classifier\_head\_accuracy: 0.7429 - val\_classifier\_head\_loss: 2.1975 - val\_loss: 2.2191 - val\_regressor\_head\_loss: 0.0215 - val\_regressor\_head\_mse: 0.0215

Epoch 96/100

**187/187** ————— 42s 226ms/step - classifier\_head\_accuracy: 0.9921 - classifier\_head\_loss: 0.0184 - loss: 0.0339 - regressor\_head\_loss: 0.0155 - regressor\_head\_mse: 0.0155 - val\_classifier\_head\_accuracy: 0.8229 - val\_classifier\_head\_loss: 1.6465 - val\_loss: 1.6665 - val\_regressor\_head\_loss: 0.0200 - val\_regressor\_head\_mse: 0.0200

Epoch 97/100

**187/187** ————— 42s 224ms/step - classifier\_head\_accuracy: 0.9860 - classifier\_head\_loss: 0.0343 - loss: 0.0502 - regressor\_head\_loss: 0.0159 - regressor\_head\_mse: 0.0159 - val\_classifier\_head\_accuracy: 0.7657 - val\_classifier\_head\_loss: 1.7489 - val\_loss: 1.7703 - val\_regressor\_head\_loss: 0.0214 - val\_regressor\_head\_mse: 0.0214

Epoch 98/100

**187/187** ————— 42s 226ms/step - classifier\_head\_accuracy: 0.9928 - classifier\_head\_loss: 0.0166 - loss: 0.0311 - regressor\_head\_loss: 0.0145 - regressor\_head\_mse: 0.0145 - val\_classifier\_head\_accuracy: 0.7657 - val\_classifier\_head\_loss: 1.7484 - val\_loss: 1.7706 - val\_regressor\_head\_loss: 0.0222 - val\_regressor\_head\_mse: 0.0222

Epoch 99/100


**187/187** ————— 43s 228ms/step - classifier\_head\_accuracy: 0.9908 - classifier\_head\_loss: 0.0210 - loss: 0.0353 - regressor\_head\_loss: 0.0142 - regressor\_head\_mse: 0.0142 - val\_classifier\_head\_accuracy: 0.8057 - val\_classifier\_head\_loss: 1.7498 - val\_loss: 1.7687 - val\_regressor\_head\_loss: 0.0189 - val\_regressor\_head\_mse: 0.0189

Epoch 100/100

**187/187** ————— 46s 247ms/step - classifier\_head\_accuracy: 0.9934 - classifier\_head\_loss: 0.0141 - loss: 0.0279 - regressor\_head\_loss: 0.0138 - regressor\_head\_mse: 0.0138 - val\_classifier\_head\_accuracy: 0.8171 - val\_classifier\_head\_loss: 1.4322 - val\_loss: 1.4530 - val\_regressor\_head\_loss: 0.0208 - val\_regressor\_head\_mse: 0.0208

Visualizing the training and validation accuracy of your classification model across epochs.

```
In [18]: plt.plot(history.history['classifier_head_accuracy'])
plt.plot(history.history['val_classifier_head_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



Defining function to calculate IoU.

```
In [21]: def intersection_over_union(boxA, boxB):
    xA = max(boxA[0], boxB[0])
    yA = max(boxA[1], boxB[1])
    xB = min(boxA[0] + boxA[2], boxB[0] + boxB[2])
    yB = min(boxA[1] + boxA[3], boxB[1] + boxB[3])

    interWidth = max(0, xB - xA)
    interHeight = max(0, yB - yA)
    interArea = interWidth * interHeight

    boxAArea = boxA[2] * boxA[3]
    boxBArea = boxB[2] * boxB[3]

    if boxAArea == 0 or boxBArea == 0:
        return 0.0

    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou
```

Defining function to prepare the test data set for testing the model.

```
In [19]: def tune_test_ds(dataset):
    dataset = dataset.map(format_instance, num_parallel_calls=tf.data.AUTOTUNE)
```

```

dataset = dataset.batch(1)
dataset = dataset.repeat()
return dataset

test_ds = tune_test_ds(raw_test_ds)

```

Visualizing test predictions for the model along with the calculated IoU.

```

In [22]: plt.figure(figsize=(12, 10))

test_list = list(test_ds.take(20).as_numpy_iterator())

print(len(test_list))

image, labels = test_list[0]

for i in range(len(test_list)):

    ax = plt.subplot(4, 5, i + 1)
    image, labels = test_list[i]

    predictions = model(image)

    predicted_box = predictions[1][0] * input_size
    predicted_box = tf.cast(predicted_box, tf.int32)

    predicted_label = predictions[0][0]

    image = image[0]

    actual_label = labels[0][0]
    actual_box = labels[1][0] * input_size
    actual_box = tf.cast(actual_box, tf.int32)

    image = image.astype("float") * 255.0
    image = image.astype(np.uint8)
    image_color = cv.cvtColor(image, cv.COLOR_GRAY2RGB)

    color = (255, 0, 0)
    # print box red if predicted and actual label do not match
    if (predicted_label[0] > 0.5 and actual_label[0] > 0) or (predicted_label[0] <
        color = (0, 255, 0)

    img_label = "unmasked"
    if predicted_label[0] > 0.5:
        img_label = "masked"

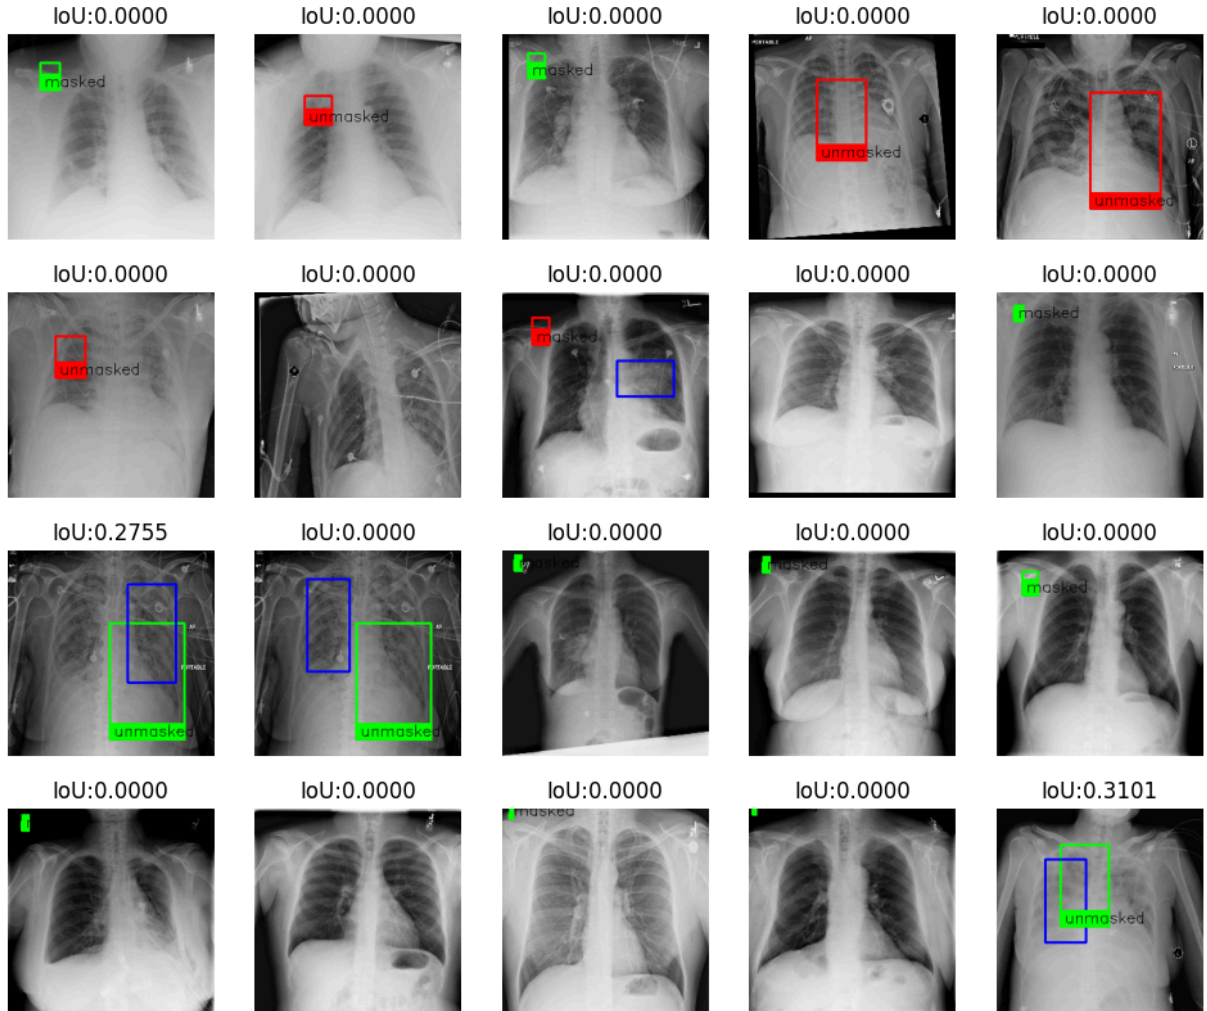
    predicted_box_n = predicted_box.numpy()
    cv.rectangle(image_color, predicted_box_n, color, 2)
    cv.rectangle(image_color, actual_box.numpy(), (0, 0, 255), 2)
    cv.rectangle(image_color, (predicted_box_n[0], predicted_box_n[1] + predicted_b
    cv.putText(image_color, img_label, (predicted_box_n[0] + 5, predicted_box_n[1]

    IoU = intersection_over_union(predicted_box.numpy(), actual_box.numpy())

```

```
plt.title("IoU:" + format(IoU, '.4f'))
plt.imshow(image_color)
plt.axis("off")
```

20



A detailed evaluation pipeline to assess your trained model's predictions on 200 test samples, visualize bounding boxes, and calculate accuracy & Intersection over Union (IoU).

```
In [23]: output_dir = "output_predictions"
os.makedirs(output_dir, exist_ok=True)

plt.figure(figsize=(12, 10))

test_list = list(test_ds.take(200).as_numpy_iterator())
print(f"Test Data Size: {len(test_list)}")

correct_count = 0
total_count = 0
iou_list = []

for i in range(len(test_list)):

    image, labels = test_list[i]
    predictions = model(image)
```

```

predicted_box = predictions[1][0] * input_size
predicted_box = tf.cast(predicted_box, tf.int32)
predicted_label = predictions[0][0]

image = image[0]
actual_label = labels[0][0]
actual_box = labels[1][0] * input_size
actual_box = tf.cast(actual_box, tf.int32)

image = image.astype("float") * 255.0
image = image.astype(np.uint8)
image_color = cv.cvtColor(image, cv.COLOR_GRAY2RGB)

color = (255, 0, 0)
if (predicted_label[0] > 0.5 and actual_label[0] > 0) or (predicted_label[0] <
    color = (0, 255, 0)
    correct_count += 1

total_count += 1

img_label = "unmasked"
if predicted_label[0] > 0.5:
    img_label = "masked"

predicted_box_n = predicted_box.numpy()
cv.rectangle(image_color, predicted_box_n, color, 2)
cv.rectangle(image_color, actual_box.numpy(), (0, 0, 255), 2)
cv.rectangle(image_color, (predicted_box_n[0], predicted_box_n[1] + predicted_b
    (predicted_box_n[0] + predicted_box_n[2], predicted_box_n[1] + pre
cv.putText(image_color, img_label, (predicted_box_n[0] + 5, predicted_box_n[1]
    cv.FONT_HERSHEY_SIMPLEX, 0.6, (0, 0, 0))

#IoU
IoU = intersection_over_union(predicted_box.numpy(), actual_box.numpy())
iou_list.append(IoU)

output_path = os.path.join(output_dir, f"prediction_{i + 1}.png")
cv.imwrite(output_path, cv.cvtColor(image_color, cv.COLOR_RGB2BGR))

accuracy = correct_count / total_count
average_iou = np.mean(iou_list)

print(f"Accuracy: {accuracy:.4f}")
print(f"Mean IoU: {average_iou:.4f}")

plt.savefig(os.path.join(output_dir, "all_predictions.png"))
plt.show()

```

Test Data Size: 200

Accuracy: 0.7550

Mean IoU: 0.0198

<Figure size 1200x1000 with 0 Axes>