

Importing relevant libraries and packages.

```
In [85]: import numpy as np
import pandas as pd
import os
import cv2 as cv
import numpy as np
from matplotlib import pyplot as plt
import pandas as pd
import pydicom
from skimage.transform import resize
import matplotlib.patches as patches
from tqdm import tqdm
import math
import graphviz
import tensorflow as tf
from tensorflow.keras.utils import plot_model
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

Load labels and display first 5 rows. Notice how the first 4 patients did not have pneumonia and do not have a box attached, while the last patient has pneumonia and has a box.

```
In [2]: train_labels = pd.read_csv('rsna-pneumonia-detection-challenge\stage_2_train_labels
train_labels.head()
```

```
Out[2]:
```

	patientId	x	y	width	height	Target
0	0004cfab-14fd-4e49-80ba-63a80b6bddd6	NaN	NaN	NaN	NaN	0
1	00313ee0-9eaa-42f4-b0ab-c148ed3241cd	NaN	NaN	NaN	NaN	0
2	00322d4d-1c29-4943-afc9-b6754be640eb	NaN	NaN	NaN	NaN	0
3	003d8fa0-6bf1-40ed-b54c-ac657f8495c5	NaN	NaN	NaN	NaN	0
4	00436515-870c-4b36-a041-de91049b9ab4	264.0	152.0	213.0	379.0	1

Creating function designed to resize an image to fit within a fixed input_size of 244 pixels while adjusting aspect ratio and bounding box (if there is one) accordingly.

```
In [3]: input_size = 244

def format_image(img, box):
    height, width = img.shape
    max_size = max(height, width)
    r = max_size / input_size
    new_width = int(width / r)
    new_height = int(height / r)
    new_size = (new_width, new_height)
    resized = cv.resize(img, new_size, interpolation=cv.INTER_LINEAR)
    new_image = np.zeros((input_size, input_size), dtype=np.uint8)
    new_image[0:new_height, 0:new_width] = resized
```

```

x, y, w, h = (box[0], box[1], box[2], box[3]) if box[0] else (0.0,0.0,0.0,0.0)
new_box = [int((x)/ r), int((y)/ r), int(w/ r), int(h/ r)] if box[0] else [0.0,

return new_image, new_box

```

Processes a sample image (00436515-...) from the training data, to visualize how the data looks. This image is one with pneumonia.

```

In [6]: datapath = 'rsna-pneumonia-detection-challenge\stage_2_train_images\00436515-870c-4
temp_img = pydicom.dcmread(datapath).pixel_array
temp_box = train_labels[train_labels['patientId'] == '00436515-870c-4b36-a041-de910


temp_img_formated, box = format_image(temp_img, temp_box)
print(box)
temp_color_img = cv.cvtColor(temp_img_formated, cv.COLOR_GRAY2RGB)

cv.rectangle(temp_color_img, box, (0, 0, 255), 1)

plt.imshow(temp_color_img)
# plt.axis("off")
plt.show()

```

[62, 36, 50, 90]



Creates and uses a function to load and process the DICOM images of Lung CTs and labels.

```

In [7]: def data_load(dataset, batch_size=3, full_data_path=r"rsna-pneumonia-detection-chal
X = []
Y = []

```

```

for index, row in tqdm(dataset.iterrows(), total=len(dataset), desc="Loading da
    filename = row['patientId']

    temp_img = pydicom.dcmread(os.path.join(full_data_path, filename + image_ex

    temp_box = [row['x'], row['y'], row['width'], row['height']] if not math.is

    img, box = format_image(temp_img, temp_box)

    img = img.astype(float) / 255.
    box = np.asarray(box, dtype=float) / input_size

    label = np.append(box, row['Target'])

    X.append(img)
    Y.append(label)

X = np.array(X)
data_X_len = len(X)
X = np.expand_dims(X, axis=3)
X = tf.convert_to_tensor(X, dtype=tf.float32)
Y = tf.convert_to_tensor(Y, dtype=tf.float32)

result = tf.data.Dataset.from_tensor_slices((X, Y))

return result, data_X_len
raw_train_ds, train_len = data_load(train_labels[:5200], ds_type="train")
print(train_len)
raw_valid_ds, valid_len = data_load(train_labels[5200:5900], ds_type="not train")
raw_test_ds, test_len = data_load(train_labels[5900:6501], ds_type="not train")

```

```

Loading data: 100%|██████████| 5200/5200 [01:00<00:00, 86.67it/s]
5200

```

```

Loading data: 100%|██████████| 700/700 [00:08<00:00, 84.23it/s]
Loading data: 100%|██████████| 601/601 [00:07<00:00, 85.53it/s]

```

Defines a function to ready the (images, label) pair to be used in TensorFlow. The two classes are "has pneumonia" and "doesn't have pneumonia," followed by the box dimensions.

```

In [8]: def format_instance(image, label):
        return image, (tf.one_hot(int(label[4]), 2), [label[0], label[1], label[2], lab

```

Defines three function to optimize the training/validation/testing dataset respectively for the tensor flow model.

```

In [10]: BATCH_SIZE = 32

def tune_training_ds(dataset):
    dataset = dataset.map(format_instance, num_parallel_calls=tf.data.AUTOTUNE)
    dataset = dataset.shuffle(1024, reshuffle_each_iteration=True)
    dataset = dataset.repeat() # The dataset be repeated indefinitely.
    dataset = dataset.batch(BATCH_SIZE)
    dataset = dataset.prefetch(tf.data.AUTOTUNE)

```

```

    return dataset

def tune_validation_ds(dataset):
    dataset = dataset.map(format_instance, num_parallel_calls=tf.data.AUTOTUNE)
    dataset = dataset.batch(len(dataset) // 4)
    dataset = dataset.repeat()
    return dataset

def tune_test_ds(dataset):
    dataset = dataset.map(format_instance, num_parallel_calls=tf.data.AUTOTUNE)
    dataset = dataset.batch(1)
    dataset = dataset.repeat()
    return dataset

```

Using the above functions, preparing the training and validation datasets.

```

In [11]: train_ds = tune_training_ds(raw_train_ds)
validation_ds = tune_validation_ds(raw_valid_ds)
test_ds = tune_test_ds(raw_test_ds)

```

Visualizing one batch of the training data set. Again, see that some patients don't have pneumonia.

```

In [12]: plt.figure(figsize=(20, 10))
for images, labels in train_ds.take(1):
    for i in range(BATCH_SIZE):
        ax = plt.subplot(4, BATCH_SIZE//4, i + 1)
        label = labels[0][i]
        box = (labels[1][i] * input_size)
        box = tf.cast(box, tf.int32)


        image = images[i].numpy().astype("float") * 255.0
        image = image.astype(np.uint8)
        image_color = cv.cvtColor(image, cv.COLOR_GRAY2RGB)

        color = (0, 0, 255)
        if label[0] > 0:
            color = (0, 255, 0)

        cv.rectangle(image_color, box.numpy(), color, 2)

    plt.imshow(image_color)
    plt.axis("off")

```



Building model architecture to perform classification (detecting pneumonia)(secondary) and bounding box regression (localizing pneumonia in X-ray images).

```
In [13]: DROPOUT_FACTOR = 0.5

def build_feature_extractor(inputs):

    x = tf.keras.layers.Conv2D(16, kernel_size=3, activation='relu', input_shape=(i
    x = tf.keras.layers.AveragePooling2D(2,2)(x)

    x = tf.keras.layers.Conv2D(32, kernel_size=3, activation = 'relu')(x)
    x = tf.keras.layers.AveragePooling2D(2,2)(x)

    x = tf.keras.layers.Conv2D(64, kernel_size=3, activation = 'relu')(x)
    x = tf.keras.layers.Dropout(DROPOUT_FACTOR)(x)
    x = tf.keras.layers.AveragePooling2D(2,2)(x)

    return x

def build_model_adaptor(inputs):
    x = tf.keras.layers.Flatten()(inputs)
    x = tf.keras.layers.Dense(64, activation='relu')(x)
    return x

def build_classifier_head(inputs):
    return tf.keras.layers.Dense(2, activation='softmax', name = 'classifier_head')

def build_regressor_head(inputs):
    return tf.keras.layers.Dense(units = 4, name = 'regressor_head')(inputs)

def build_model(inputs):

    feature_extractor = build_feature_extractor(inputs)

    model_adaptor = build_model_adaptor(feature_extractor)
```

```

classification_head = build_classifier_head(model_adaptor)

regressor_head = build_regressor_head(model_adaptor)

model = tf.keras.Model(inputs = inputs, outputs = [classification_head, regressor_head])

model.compile(optimizer=tf.keras.optimizers.Adam(),
              loss = {'classifier_head' : 'categorical_crossentropy', 'regressor_head' : 'mse'},
              metrics = {'classifier_head' : 'accuracy', 'regressor_head' : 'mse' })

return model

```

Initializing the model with the standardized image dimensions and displaying the model structure.

```

In [14]: model = build_model(tf.keras.layers.Input(shape=(input_size, input_size, 1)))

model.summary()

```

C:\Users\Sanan\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfr
a8p0\LocalCache\local-packages\Python311\site-packages\keras\src\layers\convolutional
base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to
a layer. When using Sequential models, prefer using an `Input(shape)` object as the
first layer in the model instead.

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

Model: "functional"

Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 244, 244, 1)	0	-
conv2d (Conv2D)	(None, 242, 242, 16)	160	input_layer[0][0]
average_pooling2d (AveragePooling2D)	(None, 121, 121, 16)	0	conv2d[0][0]
conv2d_1 (Conv2D)	(None, 119, 119, 32)	4,640	average_pooling2...
average_pooling2d_1 (AveragePooling2D)	(None, 59, 59, 32)	0	conv2d_1[0][0]
conv2d_2 (Conv2D)	(None, 57, 57, 64)	18,496	average_pooling2...
dropout (Dropout)	(None, 57, 57, 64)	0	conv2d_2[0][0]
average_pooling2d_2 (AveragePooling2D)	(None, 28, 28, 64)	0	dropout[0][0]
flatten (Flatten)	(None, 50176)	0	average_pooling2...
dense (Dense)	(None, 64)	3,211,328	flatten[0][0]
classifier_head (Dense)	(None, 2)	130	dense[0][0]
regressor_head (Dense)	(None, 4)	260	dense[0][0]

Total params: 3,235,014 (12.34 MB)

Trainable params: 3,235,014 (12.34 MB)

Non-trainable params: 0 (0.00 B)

Visualizing model strucutre in another way.

```
In [16]: plot_model(model, show_shapes=True, show_layer_names=True)
```

You must install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) for `plot_model` to work.

Training the model with 100 epochs.

```
In [17]: history = model.fit(train_ds,
                             steps_per_epoch=(len(raw_train_ds) // BATCH_SIZE),
                             validation_data=validation_ds, validation_steps=1,
                             epochs=100)
```

Epoch 1/100

162/162 ————— **40s** 217ms/step - classifier_head_accuracy: 0.6452 - classifier_head_loss: 0.6567 - loss: 0.7293 - regressor_head_loss: 0.0725 - regressor_head_mse: 0.0725 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 0.3978 - val_loss: 0.4143 - val_regressor_head_loss: 0.0164 - val_regressor_head_mse: 0.0164

Epoch 2/100

162/162 ————— **35s** 214ms/step - classifier_head_accuracy: 0.7615 - classifier_head_loss: 0.5063 - loss: 0.5366 - regressor_head_loss: 0.0303 - regressor_head_mse: 0.0303 - val_classifier_head_accuracy: 0.7657 - val_classifier_head_loss: 0.4577 - val_loss: 0.4764 - val_regressor_head_loss: 0.0186 - val_regressor_head_mse: 0.0186

Epoch 3/100

162/162 ————— **34s** 208ms/step - classifier_head_accuracy: 0.7741 - classifier_head_loss: 0.4907 - loss: 0.5208 - regressor_head_loss: 0.0300 - regressor_head_mse: 0.0300 - val_classifier_head_accuracy: 0.7600 - val_classifier_head_loss: 0.4983 - val_loss: 0.5171 - val_regressor_head_loss: 0.0188 - val_regressor_head_mse: 0.0188

Epoch 4/100

162/162 ————— **34s** 207ms/step - classifier_head_accuracy: 0.7893 - classifier_head_loss: 0.4619 - loss: 0.4888 - regressor_head_loss: 0.0269 - regressor_head_mse: 0.0269 - val_classifier_head_accuracy: 0.8286 - val_classifier_head_loss: 0.4157 - val_loss: 0.4302 - val_regressor_head_loss: 0.0145 - val_regressor_head_mse: 0.0145

Epoch 5/100

162/162 ————— **34s** 207ms/step - classifier_head_accuracy: 0.7894 - classifier_head_loss: 0.4507 - loss: 0.4774 - regressor_head_loss: 0.0267 - regressor_head_mse: 0.0267 - val_classifier_head_accuracy: 0.7771 - val_classifier_head_loss: 0.4703 - val_loss: 0.4874 - val_regressor_head_loss: 0.0171 - val_regressor_head_mse: 0.0171

Epoch 6/100

162/162 ————— **34s** 208ms/step - classifier_head_accuracy: 0.8094 - classifier_head_loss: 0.4247 - loss: 0.4506 - regressor_head_loss: 0.0259 - regressor_head_mse: 0.0259 - val_classifier_head_accuracy: 0.7886 - val_classifier_head_loss: 0.4398 - val_loss: 0.4601 - val_regressor_head_loss: 0.0203 - val_regressor_head_mse: 0.0203

Epoch 7/100

162/162 ————— **34s** 210ms/step - classifier_head_accuracy: 0.8139 - classifier_head_loss: 0.4126 - loss: 0.4397 - regressor_head_loss: 0.0272 - regressor_head_mse: 0.0272 - val_classifier_head_accuracy: 0.7371 - val_classifier_head_loss: 0.5485 - val_loss: 0.5710 - val_regressor_head_loss: 0.0225 - val_regressor_head_mse: 0.0225

Epoch 8/100

162/162 ————— **34s** 209ms/step - classifier_head_accuracy: 0.8354 - classifier_head_loss: 0.3653 - loss: 0.3910 - regressor_head_loss: 0.0256 - regressor_head_mse: 0.0256 - val_classifier_head_accuracy: 0.7600 - val_classifier_head_loss: 0.5039 - val_loss: 0.5304 - val_regressor_head_loss: 0.0266 - val_regressor_head_mse: 0.0266

Epoch 9/100

162/162 ————— **34s** 213ms/step - classifier_head_accuracy: 0.8473 - classifier_head_loss: 0.3372 - loss: 0.3624 - regressor_head_loss: 0.0252 - regressor_head_mse: 0.0252 - val_classifier_head_accuracy: 0.7714 - val_classifier_head_loss: 0.5124 - val_loss: 0.5351 - val_regressor_head_loss: 0.0226 - val_regressor_head_mse: 0.0226

Epoch 10/100

162/162 ————— **37s** 226ms/step - classifier_head_accuracy: 0.8663 - cla

ssifier_head_loss: 0.3098 - loss: 0.3344 - regressor_head_loss: 0.0246 - regressor_head_mse: 0.0246 - val_classifier_head_accuracy: 0.7600 - val_classifier_head_loss: 0.5920 - val_loss: 0.6182 - val_regressor_head_loss: 0.0262 - val_regressor_head_mse: 0.0262

Epoch 11/100

162/162 ————— 37s 227ms/step - classifier_head_accuracy: 0.8817 - classifier_head_loss: 0.2754 - loss: 0.3009 - regressor_head_loss: 0.0255 - regressor_head_mse: 0.0255 - val_classifier_head_accuracy: 0.7371 - val_classifier_head_loss: 0.6680 - val_loss: 0.6967 - val_regressor_head_loss: 0.0287 - val_regressor_head_mse: 0.0287

Epoch 12/100

162/162 ————— 34s 209ms/step - classifier_head_accuracy: 0.9083 - classifier_head_loss: 0.2265 - loss: 0.2495 - regressor_head_loss: 0.0230 - regressor_head_mse: 0.0230 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 0.5077 - val_loss: 0.5315 - val_regressor_head_loss: 0.0238 - val_regressor_head_mse: 0.0238

Epoch 13/100

162/162 ————— 34s 213ms/step - classifier_head_accuracy: 0.9247 - classifier_head_loss: 0.1963 - loss: 0.2175 - regressor_head_loss: 0.0212 - regressor_head_mse: 0.0212 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 0.5909 - val_loss: 0.6187 - val_regressor_head_loss: 0.0278 - val_regressor_head_mse: 0.0278

Epoch 14/100

162/162 ————— 35s 215ms/step - classifier_head_accuracy: 0.9446 - classifier_head_loss: 0.1424 - loss: 0.1637 - regressor_head_loss: 0.0213 - regressor_head_mse: 0.0213 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 0.7058 - val_loss: 0.7315 - val_regressor_head_loss: 0.0256 - val_regressor_head_mse: 0.0256

Epoch 15/100

162/162 ————— 35s 213ms/step - classifier_head_accuracy: 0.9529 - classifier_head_loss: 0.1195 - loss: 0.1391 - regressor_head_loss: 0.0196 - regressor_head_mse: 0.0196 - val_classifier_head_accuracy: 0.7371 - val_classifier_head_loss: 0.8945 - val_loss: 0.9232 - val_regressor_head_loss: 0.0287 - val_regressor_head_mse: 0.0287

Epoch 16/100

162/162 ————— 37s 225ms/step - classifier_head_accuracy: 0.9671 - classifier_head_loss: 0.0898 - loss: 0.1091 - regressor_head_loss: 0.0192 - regressor_head_mse: 0.0192 - val_classifier_head_accuracy: 0.7829 - val_classifier_head_loss: 0.6903 - val_loss: 0.7120 - val_regressor_head_loss: 0.0217 - val_regressor_head_mse: 0.0217

Epoch 17/100

162/162 ————— 36s 220ms/step - classifier_head_accuracy: 0.9672 - classifier_head_loss: 0.0903 - loss: 0.1082 - regressor_head_loss: 0.0179 - regressor_head_mse: 0.0179 - val_classifier_head_accuracy: 0.8229 - val_classifier_head_loss: 0.7706 - val_loss: 0.7938 - val_regressor_head_loss: 0.0232 - val_regressor_head_mse: 0.0232

Epoch 18/100

162/162 ————— 35s 216ms/step - classifier_head_accuracy: 0.9687 - classifier_head_loss: 0.0839 - loss: 0.1005 - regressor_head_loss: 0.0166 - regressor_head_mse: 0.0166 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 0.7477 - val_loss: 0.7710 - val_regressor_head_loss: 0.0233 - val_regressor_head_mse: 0.0233

Epoch 19/100

162/162 ————— 34s 209ms/step - classifier_head_accuracy: 0.9849 - classifier_head_loss: 0.0482 - loss: 0.0633 - regressor_head_loss: 0.0151 - regressor_head_mse: 0.0151 - val_classifier_head_accuracy: 0.8000 - val_classifier_head_loss:

1.0263 - val_loss: 1.0482 - val_regressor_head_loss: 0.0220 - val_regressor_head_mse: 0.0220
Epoch 20/100
162/162 ————— 35s 214ms/step - classifier_head_accuracy: 0.9860 - classifier_head_loss: 0.0406 - loss: 0.0551 - regressor_head_loss: 0.0145 - regressor_head_mse: 0.0145 - val_classifier_head_accuracy: 0.7714 - val_classifier_head_loss: 1.1319 - val_loss: 1.1519 - val_regressor_head_loss: 0.0200 - val_regressor_head_mse: 0.0200
Epoch 21/100
162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9931 - classifier_head_loss: 0.0231 - loss: 0.0371 - regressor_head_loss: 0.0139 - regressor_head_mse: 0.0139 - val_classifier_head_accuracy: 0.7829 - val_classifier_head_loss: 1.1090 - val_loss: 1.1299 - val_regressor_head_loss: 0.0208 - val_regressor_head_mse: 0.0208
Epoch 22/100
162/162 ————— 34s 210ms/step - classifier_head_accuracy: 0.9887 - classifier_head_loss: 0.0372 - loss: 0.0526 - regressor_head_loss: 0.0153 - regressor_head_mse: 0.0153 - val_classifier_head_accuracy: 0.7943 - val_classifier_head_loss: 1.1427 - val_loss: 1.1616 - val_regressor_head_loss: 0.0190 - val_regressor_head_mse: 0.0190
Epoch 23/100
162/162 ————— 34s 212ms/step - classifier_head_accuracy: 0.9950 - classifier_head_loss: 0.0220 - loss: 0.0353 - regressor_head_loss: 0.0133 - regressor_head_mse: 0.0133 - val_classifier_head_accuracy: 0.7829 - val_classifier_head_loss: 1.3353 - val_loss: 1.3548 - val_regressor_head_loss: 0.0195 - val_regressor_head_mse: 0.0195
Epoch 24/100
162/162 ————— 35s 214ms/step - classifier_head_accuracy: 0.9943 - classifier_head_loss: 0.0232 - loss: 0.0368 - regressor_head_loss: 0.0136 - regressor_head_mse: 0.0136 - val_classifier_head_accuracy: 0.7829 - val_classifier_head_loss: 1.0506 - val_loss: 1.0711 - val_regressor_head_loss: 0.0205 - val_regressor_head_mse: 0.0205
Epoch 25/100
162/162 ————— 34s 213ms/step - classifier_head_accuracy: 0.9940 - classifier_head_loss: 0.0219 - loss: 0.0354 - regressor_head_loss: 0.0135 - regressor_head_mse: 0.0135 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.1778 - val_loss: 1.1978 - val_regressor_head_loss: 0.0200 - val_regressor_head_mse: 0.0200
Epoch 26/100
162/162 ————— 34s 209ms/step - classifier_head_accuracy: 0.9959 - classifier_head_loss: 0.0136 - loss: 0.0256 - regressor_head_loss: 0.0120 - regressor_head_mse: 0.0120 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.0799 - val_loss: 1.0989 - val_regressor_head_loss: 0.0190 - val_regressor_head_mse: 0.0190
Epoch 27/100
162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9977 - classifier_head_loss: 0.0133 - loss: 0.0256 - regressor_head_loss: 0.0124 - regressor_head_mse: 0.0124 - val_classifier_head_accuracy: 0.8229 - val_classifier_head_loss: 1.2098 - val_loss: 1.2291 - val_regressor_head_loss: 0.0193 - val_regressor_head_mse: 0.0193
Epoch 28/100
162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9992 - classifier_head_loss: 0.0052 - loss: 0.0161 - regressor_head_loss: 0.0109 - regressor_head_mse: 0.0109 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.3813 - val_loss: 1.4012 - val_regressor_head_loss: 0.0199 - val_regressor_head_mse: 0.0199

Epoch 29/100

162/162 ————— **34s** 210ms/step - classifier_head_accuracy: 0.9969 - classifier_head_loss: 0.0122 - loss: 0.0243 - regressor_head_loss: 0.0121 - regressor_head_mse: 0.0121 - val_classifier_head_accuracy: 0.8229 - val_classifier_head_loss: 1.0791 - val_loss: 1.0977 - val_regressor_head_loss: 0.0186 - val_regressor_head_mse: 0.0186

Epoch 30/100

162/162 ————— **34s** 212ms/step - classifier_head_accuracy: 0.9937 - classifier_head_loss: 0.0213 - loss: 0.0335 - regressor_head_loss: 0.0122 - regressor_head_mse: 0.0122 - val_classifier_head_accuracy: 0.7771 - val_classifier_head_loss: 1.2198 - val_loss: 1.2407 - val_regressor_head_loss: 0.0209 - val_regressor_head_mse: 0.0209

Epoch 31/100

162/162 ————— **34s** 211ms/step - classifier_head_accuracy: 0.9974 - classifier_head_loss: 0.0109 - loss: 0.0232 - regressor_head_loss: 0.0123 - regressor_head_mse: 0.0123 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.5600 - val_loss: 1.5830 - val_regressor_head_loss: 0.0230 - val_regressor_head_mse: 0.0230

Epoch 32/100

162/162 ————— **34s** 211ms/step - classifier_head_accuracy: 0.9966 - classifier_head_loss: 0.0105 - loss: 0.0222 - regressor_head_loss: 0.0117 - regressor_head_mse: 0.0117 - val_classifier_head_accuracy: 0.7657 - val_classifier_head_loss: 1.6427 - val_loss: 1.6646 - val_regressor_head_loss: 0.0219 - val_regressor_head_mse: 0.0219

Epoch 33/100

162/162 ————— **34s** 211ms/step - classifier_head_accuracy: 0.9978 - classifier_head_loss: 0.0074 - loss: 0.0188 - regressor_head_loss: 0.0114 - regressor_head_mse: 0.0114 - val_classifier_head_accuracy: 0.7943 - val_classifier_head_loss: 1.2235 - val_loss: 1.2445 - val_regressor_head_loss: 0.0210 - val_regressor_head_mse: 0.0210

Epoch 34/100

162/162 ————— **34s** 211ms/step - classifier_head_accuracy: 0.9961 - classifier_head_loss: 0.0145 - loss: 0.0254 - regressor_head_loss: 0.0109 - regressor_head_mse: 0.0109 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.5171 - val_loss: 1.5365 - val_regressor_head_loss: 0.0194 - val_regressor_head_mse: 0.0194

Epoch 35/100

162/162 ————— **34s** 211ms/step - classifier_head_accuracy: 0.9998 - classifier_head_loss: 0.0021 - loss: 0.0126 - regressor_head_loss: 0.0105 - regressor_head_mse: 0.0105 - val_classifier_head_accuracy: 0.7943 - val_classifier_head_loss: 1.6815 - val_loss: 1.7016 - val_regressor_head_loss: 0.0202 - val_regressor_head_mse: 0.0202

Epoch 36/100

162/162 ————— **34s** 211ms/step - classifier_head_accuracy: 0.9987 - classifier_head_loss: 0.0029 - loss: 0.0127 - regressor_head_loss: 0.0098 - regressor_head_mse: 0.0098 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.4336 - val_loss: 1.4541 - val_regressor_head_loss: 0.0205 - val_regressor_head_mse: 0.0205

Epoch 37/100

162/162 ————— **35s** 215ms/step - classifier_head_accuracy: 0.9960 - classifier_head_loss: 0.0141 - loss: 0.0256 - regressor_head_loss: 0.0115 - regressor_head_mse: 0.0115 - val_classifier_head_accuracy: 0.7714 - val_classifier_head_loss: 1.2205 - val_loss: 1.2437 - val_regressor_head_loss: 0.0232 - val_regressor_head_mse: 0.0232

Epoch 38/100

162/162 ————— **34s** 211ms/step - classifier_head_accuracy: 0.9947 - cla

ssifier_head_loss: 0.0131 - loss: 0.0255 - regressor_head_loss: 0.0124 - regressor_head_mse: 0.0124 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.0944 - val_loss: 1.1109 - val_regressor_head_loss: 0.0165 - val_regressor_head_mse: 0.0165

Epoch 39/100

162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9966 - classifier_head_loss: 0.0149 - loss: 0.0259 - regressor_head_loss: 0.0110 - regressor_head_mse: 0.0110 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.6370 - val_loss: 1.6571 - val_regressor_head_loss: 0.0201 - val_regressor_head_mse: 0.0201

Epoch 40/100

162/162 ————— 34s 213ms/step - classifier_head_accuracy: 0.9988 - classifier_head_loss: 0.0065 - loss: 0.0167 - regressor_head_loss: 0.0102 - regressor_head_mse: 0.0102 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.4408 - val_loss: 1.4603 - val_regressor_head_loss: 0.0195 - val_regressor_head_mse: 0.0195

Epoch 41/100

162/162 ————— 34s 210ms/step - classifier_head_accuracy: 0.9980 - classifier_head_loss: 0.0052 - loss: 0.0150 - regressor_head_loss: 0.0098 - regressor_head_mse: 0.0098 - val_classifier_head_accuracy: 0.8343 - val_classifier_head_loss: 1.6916 - val_loss: 1.7093 - val_regressor_head_loss: 0.0177 - val_regressor_head_mse: 0.0177

Epoch 42/100

162/162 ————— 34s 212ms/step - classifier_head_accuracy: 0.9980 - classifier_head_loss: 0.0061 - loss: 0.0165 - regressor_head_loss: 0.0104 - regressor_head_mse: 0.0104 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.4380 - val_loss: 1.4579 - val_regressor_head_loss: 0.0199 - val_regressor_head_mse: 0.0199

Epoch 43/100

162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9992 - classifier_head_loss: 0.0031 - loss: 0.0131 - regressor_head_loss: 0.0100 - regressor_head_mse: 0.0100 - val_classifier_head_accuracy: 0.8343 - val_classifier_head_loss: 1.4816 - val_loss: 1.5005 - val_regressor_head_loss: 0.0189 - val_regressor_head_mse: 0.0189

Epoch 44/100

162/162 ————— 34s 212ms/step - classifier_head_accuracy: 0.9990 - classifier_head_loss: 0.0028 - loss: 0.0123 - regressor_head_loss: 0.0096 - regressor_head_mse: 0.0096 - val_classifier_head_accuracy: 0.8286 - val_classifier_head_loss: 1.4306 - val_loss: 1.4506 - val_regressor_head_loss: 0.0200 - val_regressor_head_mse: 0.0200

Epoch 45/100

162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9983 - classifier_head_loss: 0.0063 - loss: 0.0164 - regressor_head_loss: 0.0101 - regressor_head_mse: 0.0101 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.6602 - val_loss: 1.6809 - val_regressor_head_loss: 0.0207 - val_regressor_head_mse: 0.0207

Epoch 46/100

162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9994 - classifier_head_loss: 0.0036 - loss: 0.0137 - regressor_head_loss: 0.0101 - regressor_head_mse: 0.0101 - val_classifier_head_accuracy: 0.8000 - val_classifier_head_loss: 1.8120 - val_loss: 1.8312 - val_regressor_head_loss: 0.0192 - val_regressor_head_mse: 0.0192

Epoch 47/100

162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9993 - classifier_head_loss: 0.0023 - loss: 0.0116 - regressor_head_loss: 0.0093 - regressor_head_mse: 0.0093 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss:

1.7646 - val_loss: 1.7830 - val_regressor_head_loss: 0.0183 - val_regressor_head_mse: 0.0183
Epoch 48/100
162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9996 - classifier_head_loss: 0.0015 - loss: 0.0106 - regressor_head_loss: 0.0091 - regressor_head_mse: 0.0091 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.5256 - val_loss: 1.5426 - val_regressor_head_loss: 0.0171 - val_regressor_head_mse: 0.0171
Epoch 49/100
162/162 ————— 35s 214ms/step - classifier_head_accuracy: 0.9986 - classifier_head_loss: 0.0038 - loss: 0.0135 - regressor_head_loss: 0.0097 - regressor_head_mse: 0.0097 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.5835 - val_loss: 1.6023 - val_regressor_head_loss: 0.0188 - val_regressor_head_mse: 0.0188
Epoch 50/100
162/162 ————— 34s 210ms/step - classifier_head_accuracy: 0.9973 - classifier_head_loss: 0.0110 - loss: 0.0208 - regressor_head_loss: 0.0098 - regressor_head_mse: 0.0098 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.7233 - val_loss: 1.7447 - val_regressor_head_loss: 0.0213 - val_regressor_head_mse: 0.0213
Epoch 51/100
162/162 ————— 34s 212ms/step - classifier_head_accuracy: 0.9990 - classifier_head_loss: 0.0037 - loss: 0.0135 - regressor_head_loss: 0.0098 - regressor_head_mse: 0.0098 - val_classifier_head_accuracy: 0.8286 - val_classifier_head_loss: 1.3194 - val_loss: 1.3351 - val_regressor_head_loss: 0.0157 - val_regressor_head_mse: 0.0157
Epoch 52/100
162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9986 - classifier_head_loss: 0.0048 - loss: 0.0143 - regressor_head_loss: 0.0095 - regressor_head_mse: 0.0095 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.4620 - val_loss: 1.4825 - val_regressor_head_loss: 0.0205 - val_regressor_head_mse: 0.0205
Epoch 53/100
162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9980 - classifier_head_loss: 0.0048 - loss: 0.0143 - regressor_head_loss: 0.0095 - regressor_head_mse: 0.0095 - val_classifier_head_accuracy: 0.8286 - val_classifier_head_loss: 1.5368 - val_loss: 1.5559 - val_regressor_head_loss: 0.0191 - val_regressor_head_mse: 0.0191
Epoch 54/100
162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9998 - classifier_head_loss: 8.2562e-04 - loss: 0.0092 - regressor_head_loss: 0.0084 - regressor_head_mse: 0.0084 - val_classifier_head_accuracy: 0.8000 - val_classifier_head_loss: 1.5467 - val_loss: 1.5653 - val_regressor_head_loss: 0.0186 - val_regressor_head_mse: 0.0186
Epoch 55/100
162/162 ————— 34s 213ms/step - classifier_head_accuracy: 0.9977 - classifier_head_loss: 0.0071 - loss: 0.0164 - regressor_head_loss: 0.0093 - regressor_head_mse: 0.0093 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.4298 - val_loss: 1.4514 - val_regressor_head_loss: 0.0216 - val_regressor_head_mse: 0.0216
Epoch 56/100
162/162 ————— 34s 212ms/step - classifier_head_accuracy: 0.9969 - classifier_head_loss: 0.0131 - loss: 0.0239 - regressor_head_loss: 0.0107 - regressor_head_mse: 0.0107 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.2084 - val_loss: 1.2268 - val_regressor_head_loss: 0.0184 - val_regressor_head_mse: 0.0184

Epoch 57/100

162/162 ————— 35s 216ms/step - classifier_head_accuracy: 0.9983 - classifier_head_loss: 0.0069 - loss: 0.0169 - regressor_head_loss: 0.0100 - regressor_head_mse: 0.0100 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.3518 - val_loss: 1.3701 - val_regressor_head_loss: 0.0183 - val_regressor_head_mse: 0.0183

Epoch 58/100

162/162 ————— 34s 212ms/step - classifier_head_accuracy: 0.9992 - classifier_head_loss: 0.0020 - loss: 0.0113 - regressor_head_loss: 0.0093 - regressor_head_mse: 0.0093 - val_classifier_head_accuracy: 0.7943 - val_classifier_head_loss: 1.5631 - val_loss: 1.5817 - val_regressor_head_loss: 0.0185 - val_regressor_head_mse: 0.0185

Epoch 59/100

162/162 ————— 34s 212ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 4.1217e-04 - loss: 0.0092 - regressor_head_loss: 0.0088 - regressor_head_mse: 0.0088 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.4780 - val_loss: 1.4956 - val_regressor_head_loss: 0.0176 - val_regressor_head_mse: 0.0176

Epoch 60/100

162/162 ————— 34s 212ms/step - classifier_head_accuracy: 0.9999 - classifier_head_loss: 5.5547e-04 - loss: 0.0085 - regressor_head_loss: 0.0080 - regressor_head_mse: 0.0080 - val_classifier_head_accuracy: 0.8229 - val_classifier_head_loss: 1.4814 - val_loss: 1.4993 - val_regressor_head_loss: 0.0179 - val_regressor_head_mse: 0.0179

Epoch 61/100

162/162 ————— 34s 211ms/step - classifier_head_accuracy: 0.9989 - classifier_head_loss: 0.0013 - loss: 0.0093 - regressor_head_loss: 0.0080 - regressor_head_mse: 0.0080 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.7696 - val_loss: 1.7881 - val_regressor_head_loss: 0.0185 - val_regressor_head_mse: 0.0185

Epoch 62/100

162/162 ————— 34s 211ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 3.4875e-04 - loss: 0.0081 - regressor_head_loss: 0.0077 - regressor_head_mse: 0.0077 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.6263 - val_loss: 1.6437 - val_regressor_head_loss: 0.0174 - val_regressor_head_mse: 0.0174

Epoch 63/100

162/162 ————— 34s 212ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 2.4859e-04 - loss: 0.0080 - regressor_head_loss: 0.0077 - regressor_head_mse: 0.0077 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.7244 - val_loss: 1.7428 - val_regressor_head_loss: 0.0184 - val_regressor_head_mse: 0.0184

Epoch 64/100

162/162 ————— 35s 214ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 2.6890e-04 - loss: 0.0077 - regressor_head_loss: 0.0074 - regressor_head_mse: 0.0074 - val_classifier_head_accuracy: 0.8000 - val_classifier_head_loss: 1.5679 - val_loss: 1.5855 - val_regressor_head_loss: 0.0177 - val_regressor_head_mse: 0.0177

Epoch 65/100

162/162 ————— 35s 216ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 1.9765e-04 - loss: 0.0079 - regressor_head_loss: 0.0077 - regressor_head_mse: 0.0077 - val_classifier_head_accuracy: 0.7886 - val_classifier_head_loss: 1.8826 - val_loss: 1.9005 - val_regressor_head_loss: 0.0179 - val_regressor_head_mse: 0.0179

Epoch 66/100

162/162 ————— 36s 223ms/step - classifier_head_accuracy: 1.0000 - cla

ssifier_head_loss: 4.2061e-04 - loss: 0.0077 - regressor_head_loss: 0.0073 - regressor_head_mse: 0.0073 - val_classifier_head_accuracy: 0.8229 - val_classifier_head_loss: 1.6940 - val_loss: 1.7121 - val_regressor_head_loss: 0.0181 - val_regressor_head_mse: 0.0181

Epoch 67/100

162/162 ————— 35s 217ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 1.1513e-04 - loss: 0.0072 - regressor_head_loss: 0.0071 - regressor_head_mse: 0.0071 - val_classifier_head_accuracy: 0.8229 - val_classifier_head_loss: 1.6397 - val_loss: 1.6567 - val_regressor_head_loss: 0.0170 - val_regressor_head_mse: 0.0170

Epoch 68/100

162/162 ————— 36s 222ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 3.3160e-04 - loss: 0.0075 - regressor_head_loss: 0.0072 - regressor_head_mse: 0.0072 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.6849 - val_loss: 1.7025 - val_regressor_head_loss: 0.0176 - val_regressor_head_mse: 0.0176

Epoch 69/100

162/162 ————— 35s 217ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 1.2766e-04 - loss: 0.0070 - regressor_head_loss: 0.0068 - regressor_head_mse: 0.0068 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.6986 - val_loss: 1.7166 - val_regressor_head_loss: 0.0180 - val_regressor_head_mse: 0.0180

Epoch 70/100

162/162 ————— 35s 216ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 2.8524e-04 - loss: 0.0074 - regressor_head_loss: 0.0071 - regressor_head_mse: 0.0071 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.6423 - val_loss: 1.6600 - val_regressor_head_loss: 0.0177 - val_regressor_head_mse: 0.0177

Epoch 71/100

162/162 ————— 35s 217ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 8.8166e-05 - loss: 0.0068 - regressor_head_loss: 0.0067 - regressor_head_mse: 0.0067 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.6678 - val_loss: 1.6851 - val_regressor_head_loss: 0.0172 - val_regressor_head_mse: 0.0172

Epoch 72/100

162/162 ————— 35s 217ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 5.8608e-04 - loss: 0.0078 - regressor_head_loss: 0.0073 - regressor_head_mse: 0.0073 - val_classifier_head_accuracy: 0.8229 - val_classifier_head_loss: 1.4933 - val_loss: 1.5087 - val_regressor_head_loss: 0.0154 - val_regressor_head_mse: 0.0154

Epoch 73/100

162/162 ————— 35s 217ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 2.2080e-04 - loss: 0.0071 - regressor_head_loss: 0.0068 - regressor_head_mse: 0.0068 - val_classifier_head_accuracy: 0.7886 - val_classifier_head_loss: 1.7092 - val_loss: 1.7277 - val_regressor_head_loss: 0.0185 - val_regressor_head_mse: 0.0185

Epoch 74/100

162/162 ————— 35s 216ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 8.9695e-05 - loss: 0.0065 - regressor_head_loss: 0.0065 - regressor_head_mse: 0.0065 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.7266 - val_loss: 1.7445 - val_regressor_head_loss: 0.0179 - val_regressor_head_mse: 0.0179

Epoch 75/100

162/162 ————— 35s 218ms/step - classifier_head_accuracy: 0.9996 - classifier_head_loss: 4.0032e-04 - loss: 0.0071 - regressor_head_loss: 0.0067 - regressor_head_mse: 0.0067 - val_classifier_head_accuracy: 0.7486 - val_classifier_head_loss:

s: 1.7179 - val_loss: 1.7371 - val_regressor_head_loss: 0.0192 - val_regressor_head_mse: 0.0192
Epoch 76/100
162/162 ————— 36s 221ms/step - classifier_head_accuracy: 0.9897 - classifier_head_loss: 0.0293 - loss: 0.0402 - regressor_head_loss: 0.0108 - regressor_head_mse: 0.0108 - val_classifier_head_accuracy: 0.8000 - val_classifier_head_loss: 1.4712 - val_loss: 1.4933 - val_regressor_head_loss: 0.0222 - val_regressor_head_mse: 0.0222
Epoch 77/100
162/162 ————— 35s 219ms/step - classifier_head_accuracy: 0.9925 - classifier_head_loss: 0.0226 - loss: 0.0344 - regressor_head_loss: 0.0118 - regressor_head_mse: 0.0118 - val_classifier_head_accuracy: 0.8171 - val_classifier_head_loss: 1.1496 - val_loss: 1.1692 - val_regressor_head_loss: 0.0196 - val_regressor_head_mse: 0.0196
Epoch 78/100
162/162 ————— 36s 219ms/step - classifier_head_accuracy: 0.9968 - classifier_head_loss: 0.0111 - loss: 0.0214 - regressor_head_loss: 0.0103 - regressor_head_mse: 0.0103 - val_classifier_head_accuracy: 0.8000 - val_classifier_head_loss: 1.8183 - val_loss: 1.8365 - val_regressor_head_loss: 0.0181 - val_regressor_head_mse: 0.0181
Epoch 79/100
162/162 ————— 36s 222ms/step - classifier_head_accuracy: 0.9994 - classifier_head_loss: 0.0027 - loss: 0.0121 - regressor_head_loss: 0.0093 - regressor_head_mse: 0.0093 - val_classifier_head_accuracy: 0.7886 - val_classifier_head_loss: 1.6601 - val_loss: 1.6779 - val_regressor_head_loss: 0.0178 - val_regressor_head_mse: 0.0178
Epoch 80/100
162/162 ————— 35s 217ms/step - classifier_head_accuracy: 0.9999 - classifier_head_loss: 7.0734e-04 - loss: 0.0086 - regressor_head_loss: 0.0079 - regressor_head_mse: 0.0079 - val_classifier_head_accuracy: 0.7943 - val_classifier_head_loss: 1.7285 - val_loss: 1.7471 - val_regressor_head_loss: 0.0186 - val_regressor_head_mse: 0.0186
Epoch 81/100
162/162 ————— 36s 220ms/step - classifier_head_accuracy: 0.9986 - classifier_head_loss: 0.0037 - loss: 0.0122 - regressor_head_loss: 0.0085 - regressor_head_mse: 0.0085 - val_classifier_head_accuracy: 0.8343 - val_classifier_head_loss: 1.6328 - val_loss: 1.6500 - val_regressor_head_loss: 0.0172 - val_regressor_head_mse: 0.0172
Epoch 82/100
162/162 ————— 36s 223ms/step - classifier_head_accuracy: 0.9986 - classifier_head_loss: 0.0049 - loss: 0.0142 - regressor_head_loss: 0.0093 - regressor_head_mse: 0.0093 - val_classifier_head_accuracy: 0.7886 - val_classifier_head_loss: 1.7291 - val_loss: 1.7478 - val_regressor_head_loss: 0.0187 - val_regressor_head_mse: 0.0187
Epoch 83/100
162/162 ————— 36s 222ms/step - classifier_head_accuracy: 0.9989 - classifier_head_loss: 0.0020 - loss: 0.0106 - regressor_head_loss: 0.0086 - regressor_head_mse: 0.0086 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.7481 - val_loss: 1.7658 - val_regressor_head_loss: 0.0177 - val_regressor_head_mse: 0.0177
Epoch 84/100
162/162 ————— 36s 220ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 2.5190e-04 - loss: 0.0082 - regressor_head_loss: 0.0079 - regressor_head_mse: 0.0079 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.9020 - val_loss: 1.9218 - val_regressor_head_loss: 0.0198 - val_regressor_head_mse: 0.0198

Epoch 85/100

162/162 ————— 36s 221ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 9.9876e-05 - loss: 0.0076 - regressor_head_loss: 0.0075 - regressor_head_mse: 0.0075 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.8471 - val_loss: 1.8648 - val_regressor_head_loss: 0.0176 - val_regressor_head_mse: 0.0176

Epoch 86/100

162/162 ————— 35s 219ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 5.9785e-04 - loss: 0.0080 - regressor_head_loss: 0.0075 - regressor_head_mse: 0.0075 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.8178 - val_loss: 1.8357 - val_regressor_head_loss: 0.0179 - val_regressor_head_mse: 0.0179

Epoch 87/100

162/162 ————— 36s 219ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 3.1416e-04 - loss: 0.0074 - regressor_head_loss: 0.0071 - regressor_head_mse: 0.0071 - val_classifier_head_accuracy: 0.8000 - val_classifier_head_loss: 1.8685 - val_loss: 1.8862 - val_regressor_head_loss: 0.0178 - val_regressor_head_mse: 0.0178

Epoch 88/100

162/162 ————— 36s 220ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 2.3685e-04 - loss: 0.0072 - regressor_head_loss: 0.0070 - regressor_head_mse: 0.0070 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.9044 - val_loss: 1.9230 - val_regressor_head_loss: 0.0185 - val_regressor_head_mse: 0.0185

Epoch 89/100

162/162 ————— 36s 219ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 6.0100e-05 - loss: 0.0069 - regressor_head_loss: 0.0068 - regressor_head_mse: 0.0068 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.9018 - val_loss: 1.9203 - val_regressor_head_loss: 0.0185 - val_regressor_head_mse: 0.0185

Epoch 90/100

162/162 ————— 36s 221ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 4.1210e-05 - loss: 0.0069 - regressor_head_loss: 0.0068 - regressor_head_mse: 0.0068 - val_classifier_head_accuracy: 0.7943 - val_classifier_head_loss: 2.0295 - val_loss: 2.0485 - val_regressor_head_loss: 0.0190 - val_regressor_head_mse: 0.0190

Epoch 91/100

162/162 ————— 36s 221ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 1.4826e-04 - loss: 0.0073 - regressor_head_loss: 0.0072 - regressor_head_mse: 0.0072 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.9025 - val_loss: 1.9209 - val_regressor_head_loss: 0.0184 - val_regressor_head_mse: 0.0184

Epoch 92/100

162/162 ————— 36s 221ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 5.2011e-04 - loss: 0.0073 - regressor_head_loss: 0.0068 - regressor_head_mse: 0.0068 - val_classifier_head_accuracy: 0.8057 - val_classifier_head_loss: 1.8884 - val_loss: 1.9058 - val_regressor_head_loss: 0.0175 - val_regressor_head_mse: 0.0175

Epoch 93/100

162/162 ————— 36s 220ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 6.1657e-05 - loss: 0.0068 - regressor_head_loss: 0.0067 - regressor_head_mse: 0.0067 - val_classifier_head_accuracy: 0.8114 - val_classifier_head_loss: 1.7995 - val_loss: 1.8171 - val_regressor_head_loss: 0.0177 - val_regressor_head_mse: 0.0177

Epoch 94/100

162/162 ————— 36s 220ms/step - classifier_head_accuracy: 1.0000 - cla

ssifier_head_loss: 5.8021e-05 - loss: 0.0067 - regressor_head_loss: 0.0067 - regressor_head_mse: 0.0067 - val_classifier_head_accuracy: 0.7771 - val_classifier_head_loss: 1.8456 - val_loss: 1.8631 - val_regressor_head_loss: 0.0174 - val_regressor_head_mse: 0.0174

Epoch 95/100

162/162 ————— 36s 222ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 5.5537e-04 - loss: 0.0074 - regressor_head_loss: 0.0069 - regressor_head_mse: 0.0069 - val_classifier_head_accuracy: 0.7829 - val_classifier_head_loss: 1.8145 - val_loss: 1.8328 - val_regressor_head_loss: 0.0183 - val_regressor_head_mse: 0.0183

Epoch 96/100

162/162 ————— 36s 221ms/step - classifier_head_accuracy: 1.0000 - classifier_head_loss: 5.9073e-05 - loss: 0.0066 - regressor_head_loss: 0.0065 - regressor_head_mse: 0.0065 - val_classifier_head_accuracy: 0.7943 - val_classifier_head_loss: 1.6198 - val_loss: 1.6365 - val_regressor_head_loss: 0.0167 - val_regressor_head_mse: 0.0167

Epoch 97/100

162/162 ————— 36s 221ms/step - classifier_head_accuracy: 0.9998 - classifier_head_loss: 3.2754e-04 - loss: 0.0071 - regressor_head_loss: 0.0067 - regressor_head_mse: 0.0067 - val_classifier_head_accuracy: 0.7771 - val_classifier_head_loss: 1.6916 - val_loss: 1.7090 - val_regressor_head_loss: 0.0174 - val_regressor_head_mse: 0.0174

Epoch 98/100

162/162 ————— 36s 221ms/step - classifier_head_accuracy: 0.9997 - classifier_head_loss: 0.0012 - loss: 0.0088 - regressor_head_loss: 0.0076 - regressor_head_mse: 0.0076 - val_classifier_head_accuracy: 0.7429 - val_classifier_head_loss: 2.2307 - val_loss: 2.2541 - val_regressor_head_loss: 0.0234 - val_regressor_head_mse: 0.0234

Epoch 99/100


162/162 ————— 36s 221ms/step - classifier_head_accuracy: 0.9980 - classifier_head_loss: 0.0071 - loss: 0.0171 - regressor_head_loss: 0.0100 - regressor_head_mse: 0.0100 - val_classifier_head_accuracy: 0.7657 - val_classifier_head_loss: 1.9907 - val_loss: 2.0117 - val_regressor_head_loss: 0.0210 - val_regressor_head_mse: 0.0210

Epoch 100/100

162/162 ————— 36s 224ms/step - classifier_head_accuracy: 0.9965 - classifier_head_loss: 0.0073 - loss: 0.0170 - regressor_head_loss: 0.0098 - regressor_head_mse: 0.0098 - val_classifier_head_accuracy: 0.7600 - val_classifier_head_loss: 2.1485 - val_loss: 2.1669 - val_regressor_head_loss: 0.0184 - val_regressor_head_mse: 0.0184

Visualizing the training and validation accuracy of our classification model across epochs.

```
In [18]: plt.plot(history.history['classifier_head_accuracy'])
plt.plot(history.history['val_classifier_head_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()
```



Defining function to calculate IoU, using the formula in the kaggle challenge overview.

```
In [19]: def intersection_over_union(boxA, boxB):
    xA = max(boxA[0], boxB[0])
    yA = max(boxA[1], boxB[1])
    xB = min(boxA[0] + boxA[2], boxB[0] + boxB[2])
    yB = min(boxA[1] + boxA[3], boxB[1] + boxB[3])

    interWidth = max(0, xB - xA)
    interHeight = max(0, yB - yA)
    interArea = interWidth * interHeight

    boxAArea = boxA[2] * boxA[3]
    boxBArea = boxB[2] * boxB[3]

    if boxAArea == 0 or boxBArea == 0:
        return 0.0

    iou = interArea / float(boxAArea + boxBArea - interArea)
    return iou
```

Visualizing 20 test predictions for the model along with the calculated IoU. Red box printed if a case was wrongly classified. Otherwise, printed a green box if correctly classified. The box label is unmasked if the model classified it as not having pneumonia and masked if it was classified as having pneumonia.

```

In [ ]: plt.figure(figsize=(12, 10))

test_list = list(test_ds.take(20).as_numpy_iterator())

image, labels = test_list[0]

for i in range(len(test_list)):

    ax = plt.subplot(4, 5, i + 1)
    image, labels = test_list[i]

    predictions = model(image)

    predicted_box = predictions[1][0] * input_size
    predicted_box = tf.cast(predicted_box, tf.int32)

    predicted_label = predictions[0][0]

    image = image[0]

    actual_label = labels[0][0]
    actual_box = labels[1][0] * input_size
    actual_box = tf.cast(actual_box, tf.int32)

    image = image.astype("float") * 255.0
    image = image.astype(np.uint8)
    image_color = cv.cvtColor(image, cv.COLOR_GRAY2RGB)

    color = (255, 0, 0)
    # print box red if predicted and actual label do not match
    if (predicted_label[0] > 0.5 and actual_label[0] > 0) or (predicted_label[0] <
        color = (0, 255, 0)

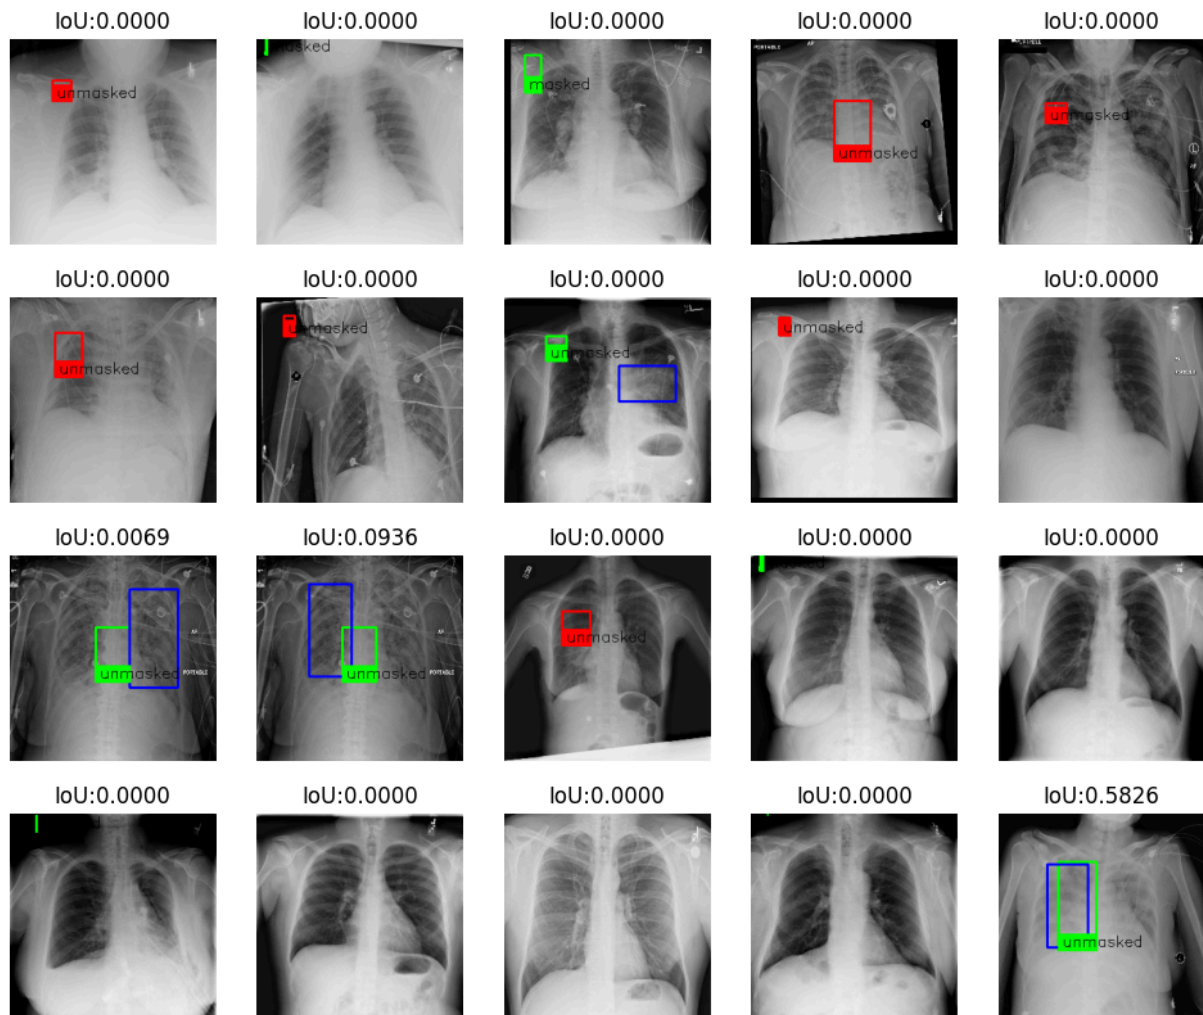
    img_label = "unmasked"
    if predicted_label[0] > 0.5:
        img_label = "masked"

    predicted_box_n = predicted_box.numpy()
    cv.rectangle(image_color, predicted_box_n, color, 2)
    cv.rectangle(image_color, actual_box.numpy(), (0, 0, 255), 2)
    cv.rectangle(image_color, (predicted_box_n[0], predicted_box_n[1] + predicted_b
    cv.putText(image_color, img_label, (predicted_box_n[0] + 5, predicted_box_n[1]

    IoU = intersection_over_union(predicted_box.numpy(), actual_box.numpy())

    plt.title("IoU:" + format(IoU, '.4f'))
    plt.imshow(image_color)
    plt.axis("off")

```



Evaluation of the trained model's predictions on test set: calculating accuracy & Intersection over Union (IoU). Also created confusion matrix for the categorization.

In [111...

```
output_dir = "output_predictions"
os.makedirs(output_dir, exist_ok=True)

plt.figure(figsize=(12, 10))

test_list = list(test_ds.take(len(raw_test_ds)).as_numpy_iterator())
print(f"Test Data Size: {len(test_list)}")

correct_count = 0
total_count = 0
iou_list = []
label_predicted = []
label_actual = []

for i in range(len(test_list)):

    image, labels = test_list[i]
    predictions = model(image)

    predicted_box = predictions[1][0] * input_size
    predicted_box = tf.cast(predicted_box, tf.int32)
```

```

predicted_label = predictions[0][0]

actual_label = labels[0][0]
actual_box = labels[1][0] * input_size
actual_box = tf.cast(actual_box, tf.int32)

if (predicted_label[0] > 0.5 and actual_label[0] > 0) or (predicted_label[0] <
    correct_count += 1

total_count += 1

img_label = "unmasked"
if predicted_label[0] > 0.5:
    img_label = "masked"

#IoU
IoU = intersection_over_union(predicted_box.numpy(), actual_box.numpy())
iou_list.append(IoU)
label_predicted.append(int(predicted_label[0]>0.5))
label_actual.append(int(actual_label[0]))

accuracy = correct_count / total_count
average_iou = np.mean(iou_list)


confusionmatrix = confusion_matrix(label_actual, label_predicted)
disp = ConfusionMatrixDisplay(confusionmatrix, display_labels=['No Pneumonia', 'Pne
disp.plot
plt.title("Confusion Matrix")
plt.show()

recall = confusionmatrix[1][1]/(confusionmatrix[1][1] + confusionmatrix[1][0])
precision = confusionmatrix[1][1]/(confusionmatrix[1][1] + confusionmatrix[0][1])
F1 = 2 * (precision * recall) / (precision + recall)
print(f"Recall: {recall:.4f}, Precision: {precision:.4f}, F1 Score: {F1:.4f}")
print(f"Accuracy: {accuracy:.4f}")
print(f"Mean IoU: {average_iou:.4f}")

```

Test Data Size: 601

<Figure size 1200x1000 with 0 Axes>



Recall: 0.7108, Precision: 0.9614, F1 Score: 0.8173

Accuracy: 0.7404

Mean IoU: 0.0148